

Model Collapse on Manifolds: a Closed-Form Capacity Floor and a Data-Anchoring Fix

Naman Chandwani
Independent Researcher
namanchandwani.10@gmail.com

July 2026

Abstract

Generative models are increasingly retrained on data that includes their own past outputs, and for diffusion models this *self-consuming* loop is known to degrade quality from one generation to the next. We settle the case that matters most in practice and that the sharpest prior theory [1] repairs only for smooth densities and explicitly leaves open—data concentrated on a thin low-dimensional manifold—with matching negative and constructive results. *Negatively*, we prove a closed-form *capacity floor*: any truncated reverse-SDE sampler whose normal score slope is capped at $\bar{\kappa}$ retains excess width of at least $1/(8\bar{\kappa}^2)$ on every pass, with no dependence on the sampling schedule. The mechanism is geometric: near a manifold of width σ the score a diffusion model must represent is steep only inside a short, finite window of noise levels, so a slope-capped network bakes in excess width there that later denoising cannot remove. A model held at a fixed slope ceiling therefore collapses without bound as the manifold thins, and staying bounded requires a slope budget growing like $\sqrt{J}$ in the Fisher information J . Throughout, the hypothesis bounds the network’s normal Lipschitz norm, not its parameter count; the link from architecture size to attainable slope is measured here, not proved. Empirically the trap is two-horned: the stochastic sampler sticks $\sim 11\times$ too wide, while the deterministic one collapses to a point; neither reaches the data. *Constructively*, we prove that periodically re-matching the model’s output to a small, *fixed* pool of real examples makes the loop *memoryless*, so that error stops compounding for every positive fraction of real data. The match rescales the output to the *local width* of the reference; it does not re-insert the reference examples into training. Over a 20-generation head-to-head the prior state of the art plateaus at $9.8\sigma^2$ while anchoring holds $1.01\sigma^2$, and on pixel-space MNIST anchoring removes $\sim 81\%$ of the off-manifold collapse. A preliminary scale-up to CIFAR-10 in raw pixel space (3,072 dimensions), across three seeds, reproduces the separation and, by three independent nearest-neighbor metrics, shows the anchor keeps full coverage of the real data, ruling out mode collapse. The core experiments run on CPU and all are reproducible from released scripts.

1 Introduction

In plain terms. Train an image generator, use it to make pictures, then train the next generator partly on those pictures, and repeat. Each round the model drifts a little further from real data, the way a photocopy of a photocopy loses detail. This is *model collapse*. Recent theory shows the drift can be repaired when the data is “smooth,” but leaves open the case that actually describes images, audio, and most real data: points that lie on a thin surface (a *manifold*) inside a much larger space. This paper handles that case. We prove there is a hard floor on how faithful a fixed-size model can

stay, set by a single number, how sharp an edge the network can represent, that no amount of tuning removes. Then we give a cheap fix: each round, gently rescale the model’s output so its fine-scale spread matches a small, fixed sample of real data. The fix stops the drift from compounding, for any positive amount of fresh real data.

Web-scale corpora already contain machine-generated data, and synthetic data is used deliberately in modern training pipelines. When a generative model is repeatedly retrained on a pool containing its own previous outputs, quality degrades across generations, a phenomenon called *model collapse* [2–5]. For diffusion models, Khelifa et al. [1] give the sharpest analysis to date: with a fraction λ of fresh real data per generation and an annealed truncation schedule $t_0(g) \rightarrow 0$, the recursion converges back to the data distribution, *provided the data has finite Fisher information* (their Case 1). They explicitly leave open their Case 2: data supported on, or concentrated near, a low-dimensional manifold, where the Fisher information $J \asymp (d - k)/\sigma^2$ diverges as the noise width $\sigma \rightarrow 0$. Essentially all natural data (images, audio, molecular conformations) is of this second kind [6, 7].

This paper answers Case 2.

The idea in one paragraph. The hardest property to preserve across generations is the data’s *thinness*—the exact width of a pen stroke, the sharp edge of a shape. A finite-capacity denoiser can reproduce that thinness only inside a brief window of noise levels, so a small, fixed amount of blur is baked in every generation and accumulates. This yields our two results, made precise below: a *negative* one—the blur has an exact floor set by a single number, the steepest slope the network can represent, that no training or sampling schedule lowers—and a *constructive* one—the accumulation is cancelled by periodically re-matching the model’s output to the *local width* of a small, unchanging set of real examples, without copying those examples back into training.

The negative half is a *capacity floor* (Section 5). The score that a diffusion model must represent near a σ -thin manifold has normal slope $\kappa(t) = s(t)/\gamma^2(t)$ which is *non-monotone* in diffusion time: it peaks at $\kappa_{\max} = 1/(2\sigma)$ near $t \approx \sigma^2$ and falls back toward zero as $t \rightarrow 0$ (Figure 1a). A network whose normal slope is bounded by $\bar{\kappa}$ therefore under-contracts on a *finite band* of diffusion times; the excess variance frozen in while crossing that band can never be removed afterwards, because too few contraction e-folds remain below the band. We prove that in the singular limit the resulting floor is a universal function of the single ratio $\rho = \bar{\kappa}/\kappa_{\max}$ and admits *closed forms*: under no assumption beyond the slope bound, the floor is exactly $\sigma^2/(2\rho^2) = 1/(8\bar{\kappa}^2)$, so a fixed network collapses relatively without bound as $\sigma \rightarrow 0$, and bounded collapse requires $\bar{\kappa} = \Omega(1/\sigma) = \Omega(\sqrt{J})$. No truncation schedule enters these statements: annealing targets a term the schedule can kill, while the floor lives in a term it cannot touch. Empirically the trap is *two-horned* (Figure 2): keeping the reverse-SDE noise, the sampler sticks $\sim 11\times$ too wide; removing it, the recursion collapses to a single point. Neither branch reaches the data.

The constructive half is an anchoring operator (Section 6). A bidirectional, per-axis, local moment match of the training pool (and sampler output) to a small *stale* reference of real data (2,000 samples drawn once at generation zero) pins the normal variance at the reference value each generation. We prove this makes the recursion *memoryless*: the fixed-point equation loses its dependence on the previous generation, the critical fraction disappears, and the dynamics of Khelifa et al. [1]’s Theorem 4.1 are restored in the singular regime. In a 20-generation head-to-head against their annealed-truncation remedy (Figure 5), the annealed arm plateaus at $9.8\sigma^2$ while the anchored arms hold $1.01\sigma^2$, flat, on both samplers we test; on real pixels (MNIST), anchoring closes $\sim 81\%$ of the collapse gap (Figure 6); and a preliminary CIFAR-10 pixel recursion (3,072 dimensions, three seeds) reproduces the anchored/unanchored split with full mode coverage (Figure 7).

Contributions.

1. **A law of collapse** (Theorem 1): under a measured one-dimensional concave response law, the variance recursion has a unique plateau for $\lambda > \lambda^* = c_\infty/(1 + c_\infty)$ and no plateau below; truncation enters only additively, which makes precise why annealing schedules cannot help singular data.
2. **A closed-form capacity floor** (Theorem 2, Corollary 2): any reverse-SDE sampler with normal score slope $\leq \bar{\kappa}$ is floored; in the singular limit the floor is $1/(8\bar{\kappa}^2)$ unconditionally and $((1 + 3e^{-4})/8)\bar{\kappa}^{-2}$ for the empirically observed no-overshoot class: two constants that differ by only 5.5%. Bounded relative collapse as $\sigma \rightarrow 0$ requires $\bar{\kappa} \geq 1/(2\sigma\sqrt{2C})$.
3. **A memorylessness theorem for anchoring** (Theorem 3): matching pool and output normal variance to a stale real reference removes the recursion’s memory, for every $\lambda > 0$, with an error bound independent of the ambient dimension.
4. **Falsifiable empirics**, each targeting a named claim:
 - *Floor anatomy & intervention*: read at the ceiling the theorem constrains, the law’s envelope explains 81% of the measured single-pass floor at $\sigma = 0.05$ (against 58% read under p_t , at a quarter of the seed spread), the remainder resolving into one further measured channel (in-band profile sag); the margin narrows with tube width, and where the flat envelope is exceeded at $\sigma = 0.13$ the curved-geometry correction of Corollary 1 removes the excess outright; and across five training protocols that move the ceiling $3.2 \rightarrow 5.8$, every measured floor stays above its proven bound.
 - *Two-horned ablation*: separates the floor’s existence (sampler noise) from its size (slope ceiling).
 - *Head-to-head*: a 20-generation, 5-seed win over the state-of-the-art remedy, with a replay-buffer control at the anchor’s exact information budget that separates anchoring from replay, and a subcritical-fraction run exhibiting the law’s no-plateau regime.
 - *Real and scaled data*: a pixel-space MNIST recursion, and a preliminary CIFAR-10 pixel recursion (three seeds) that rules out mode collapse by three independent nearest-neighbor metrics.

Throughout we keep the epistemic status of each claim explicit: *we prove* (mathematics, given stated assumptions), *we observe* (multi-seed experiments), or *we model* (measured structural laws that we do not derive). Section 8 collects everything in the third category.

2 Related work

Model collapse. Collapse under recursive training is documented and analyzed across model families [2–5, 8]; accumulation (rather than replacement) of data and sufficient fresh-data fractions mitigate it in the absolutely continuous setting. Khelifa et al. [9] bound error propagation in recursively trained diffusions via a discounted accumulation argument, but their assumptions require absolutely continuous data (structurally excluding manifolds), with fixed truncation and upper bounds rather than laws. Khelifa et al. [1] characterize the limiting collapse distribution spectrally and prove annealed truncation converges for finite-Fisher data (their Case 1), stating manifold-supported data (Case 2) as open. We give the Case 2 answer: a floor theorem for their sampler class, and an operator restoring their Case 1 dynamics.

Diffusion on manifold data. That diffusion models detect and concentrate on data manifolds is established [7, 10]. The singular ($\sigma \rightarrow 0$) score is the object of a growing one-shot literature: optimal denoising under low regularity [11], provable learning under the manifold hypothesis [12], and mitigation of the score singularity [13]. These works are one-generation; our subject is the *recursion*, where a per-generation floor compounds into a permanently displaced plateau. We use the Tweedie/posterior-mean jump [11, 14] as a raw-error reducer inside the loop and are explicit that it does not, by itself, break the floor (Section 6).

Local geometry estimation. Our anchoring operator estimates local tangent frames by k NN-PCA; its error analysis imports standard local-PCA perturbation rates [15, 16]. The operator needs only *shared* frames between pool and reference, not correct ones, which is why these rates enter only at second order (Lemma 4).

3 Setup

Data model (σ -tube). Let $M \subset \mathbb{R}^d$ be a compact smooth k -dimensional manifold with reach $\tau > 0$. Data are $X_0 = \varphi(U) + \sigma N_\perp$, with U uniform on M and $N_\perp$ standard normal in the $(d - k)$ -dimensional normal space; $\sigma > 0$ is the tube width. The singular regime is $\sigma \rightarrow 0$; the Fisher information of the marginal scales as $J \asymp (d - k)/\sigma^2 \rightarrow \infty$ (Case 2 of 1).

Forward process and the normal coordinate. We use the variance-preserving OU forward process $X_t = a(t)X_0 + s(t)Z$ with $a(t) = e^{-t/2}$, $s^2(t) = 1 - e^{-t}$ [17, 18]. Fix a point of M and work in a normal coordinate $x \in \mathbb{R}$, treating the fiber marginal as Gaussian and decoupled from tangential position (*flat-tube idealization*; curvature enters at $O(s^2/\tau)$ and is tracked where it matters). Then $X_0 \sim \mathcal{N}(0, \sigma^2)$, $X_t \sim \mathcal{N}(0, \gamma^2(t))$ with $\gamma^2 = a^2\sigma^2 + s^2$, and the denoising target is linear:

$$\varepsilon^*(x, t) = \kappa(t)x, \quad \kappa(t) = \frac{s(t)}{\gamma^2(t)}, \quad \nabla \log p_t(x) = -\frac{x}{\gamma^2(t)}. \quad (1)$$

κ is the normal slope the score network must represent. It satisfies $\kappa \approx 1/s$ for $s^2 \gg \sigma^2$, peaks at $\kappa_{\max} = 1/(2\sigma)$ near $t \approx \sigma^2$, and falls back to 0 as $t \rightarrow 0$ (Figure 1a). In the singular limit the *peak* slope diverges, which is the precise sense in which infinite Fisher information bites the estimator.

Samplers. $S_{\text{trunc}}(t_0)$ integrates the learned reverse SDE from T down to $t_0 > 0$ (Euler-Maruyama, $x \leftarrow x + (-\frac{1}{2}x + \hat{\varepsilon}/s)dt + \sqrt{|dt|}\xi$) and outputs x_{t_0} ; all prior remedies, including the annealed schedules of Khelifa et al. [1], are of this class. $S_{\text{jump}}(\hat{t})$ integrates only to a moderate $\hat{t}$ and then outputs the Tweedie posterior-mean jump $\hat{x}_0 = (x_{\hat{t}} - s(\hat{t})\hat{\varepsilon})/a(\hat{t})$ [11, 14].

The recursion. For a distribution p on $\mathbb{R}^d$ define the per-direction normal second moment $v(p) = \mathbb{E}_p[\text{dist}(X, M)^2]/(d - k)$; true data has $v = \sigma^2$. Generation g trains $\hat{\varepsilon}_g$ by denoising score matching on $\text{pool}_g = \lambda(\text{fresh real}) + (1 - \lambda)(\text{generation } g-1)$ and samples; v_g is the output value. Since v is linear in the measure, $w_g := v(\text{pool}_g) = \lambda\sigma^2 + (1 - \lambda)v_{g-1}$ exactly.

Modeling assumptions (*we model*). Our law rests on two assumptions about the trained pipeline that we *measure and cross-validate* but do not derive.

Assumption 1 (scalar reduction). *The per-generation map factors through the pool width: $v_g = F_g(w_g)$. Why a scalar: the collapse observable is a variance along the normal fibre, one number per normal axis by construction, so this is an order-parameter statement (the width closes on*

itself) rather than a claim that the network is scalar, in the same way an order parameter closes a mean-field recursion without the microscopic state being low-dimensional. Validation: fixed points measured in one experiment type predict those of an independent type within 1 to 13% across six (t_0, λ) cells (Section 7).

Assumption 2 (concave response). $F_g(w) = a^2(t_0(g))w + s^2(t_0(g)) + e^2(w)$, where the sampling excess $e^2 : [0, \infty) \rightarrow (0, \infty)$ is nondecreasing and concave with $e^2(0) = e_0^2 > 0$. In words: the output width is the pool width shrunk by the forward map (a^2w), plus the truncation remainder (s^2), plus a sampler-added excess e^2 that (i) never vanishes even for a perfectly thin pool ($e_0^2 > 0$, this is the floor), (ii) grows with a fatter pool but (iii) with diminishing returns (concave: doubling the pool width less than doubles the added excess). Why concave: the excess comes from the network failing to represent the steep part of the score, and a fatter pool is easier to represent (the required peak slope $\kappa_{\max} = 1/(2\sigma_{\text{pool}})$ falls), so each extra unit of pool width buys less extra error, exactly a saturating (concave) response. Validation: two independent estimators. The direct probe (training single-width pools and measuring $dv/dw - 1$; Appendix F) gives a local slope that is small and nonincreasing, $+0.13 \rightarrow +0.05$, across the upper half of the measured range, where all of the recursion’s fixed points sit; its small magnitude is itself a measured near-cancellation of two $O(0.7)$ channels (Section 7). The recursion’s own local convergence rates give an independent, indirect estimate that falls from $+0.6$ at the thinnest measured fixed-point pool ($\lambda = 0.75$) to ≈ 0 at the fattest ($\lambda = 0.25$): the same saturating direction, though the two estimators disagree in magnitude where they overlap, and we treat the rate-based values as directional only (rate fits must separate a transient, a known failure mode of that estimator). At the probe’s thinnest widths the slope estimates rise from 0, a regime the recursion never visits; the law’s conclusions require only a unique downcrossing of $T(v) - v$, which the measured slopes (≤ 0.13 at every fixed point, against thresholds $\lambda/(1 - \lambda) \geq 1/3$ at every plateau-forming fraction we test, $\lambda \geq 1/4$) guarantee.

4 The law of collapse

Under Assumptions 1 and 2 the recursion is one-dimensional:

$$v_g = T(v_{g-1}) + s_g^2, \quad T(v) = \lambda\sigma^2 + (1 - \lambda)v + e^2(\lambda\sigma^2 + (1 - \lambda)v), \quad (2)$$

with $s_g^2 := s^2(t_0(g))$ the truncation remainder. (The display takes $a^2(t_0) = 1$; the general case replaces s_g^2 by $\eta_g = s_g^2 + (a^2(t_0(g)) - 1)w_g$. Its second term $\rightarrow 0$ as $t_0 \rightarrow 0$ whenever w_g stays bounded (part i), and at the small fixed t_0 we use it is $O(t_0w_g)$, absorbed into the fitted law; in the divergent case (part ii) we work with the display form, where $\eta_g = s_g^2 \geq 0$ only strengthens the conclusion.) Let $c_\infty := \lim_{w \rightarrow \infty} e^{2l}(w)$ (exists by concavity) and define the *critical real-data fraction* $\lambda^* := c_\infty/(1 + c_\infty)$.

How to read the theorem. The per-generation update T pushes the width toward a balance point. Part (i): if you keep enough fresh real data ($\lambda > \lambda^*$) there is a single stable plateau the width always settles to, from any start and any schedule. Part (ii): below that fresh-data threshold there is no plateau at all, the width runs off to infinity. (For our trained pipeline the measured saturated slope puts λ^* at ≈ 0.03 – 0.05 ; Section 7 locates it and exhibits both regimes.) Part (iii): if the sampler excess happens to be a straight line in the pool width, the plateau has a clean closed form. The one thing to notice is what λ^* and v_∞ depend on: the sampler excess e^2 , *not* the truncation schedule, which is the whole point of Remark 1.

Theorem 1 (law of the collapse; we prove). Fix $\lambda \in (0, 1]$. (i) If $\lambda > \lambda^*$, T has a unique fixed point $v_\infty > 0$ solving $v = \lambda\sigma^2 + (1 - \lambda)v + e^2(\lambda\sigma^2 + (1 - \lambda)v)$, and for every $v_0 \geq 0$ and every truncation




Figure 1: **The mechanism and the law.** (a) The required normal slope $\kappa(t) = s/\gamma^2$ is non-monotone: it peaks at $\kappa_{\max} = 1/(2\sigma)$ near $t \approx \sigma^2$ and falls back to zero. A network with ceiling $\bar{\kappa} < \kappa_{\max}$ under-contracts only on a *finite band* (shaded, the cap-to-curve deficit lens); variance frozen in the band cannot be removed below it (too few contraction e-folds remain). (b) The resulting floor in units of σ^2 as a function of $\rho = 2\sigma\bar{\kappa}$: closed-form no-overshoot law $g(\rho)$ (blue, Theorem 2), exact unconditional law $1/(2\rho^2)$ (orange). Green rings: direct finite- σ ODE integration at $\sigma = 0.05$ swept across ρ (self-checked against the appendix triples), sitting just above $g(\rho)$ and following the no-overshoot branch, the relative $O(\sigma^2/\rho^2)$ convergence the theorem predicts. The dotted vertical marks $\rho = 1/\sqrt{2}$, below which the floor exceeds the true width σ^2 .

schedule with $s_g^2 \rightarrow 0$ the iterates converge, $v_g \rightarrow v_\infty$, with local rate $\rho_{\text{loc}} = (1 - \lambda)(1 + e^{2l}(w_\infty)) < 1$ (where $w_\infty = \lambda\sigma^2 + (1 - \lambda)v_\infty$ is the fixed-point pool width; the strict inequality is forced by $T(0) > 0$, Appendix A). (ii) If $(1 - \lambda)(1 + c_\infty) > 1$, then $T(v) - v \geq T(0) > 0$ for all v and $v_g \rightarrow \infty$: no plateau, for any schedule. (iii) In the affine special case $e^2(w) = e_0^2 + cw$: $\lambda^* = c/(1 + c)$, $\rho_{\text{loc}} = (1 + c)(1 - \lambda)$, and $v_\infty = ((1 + c)\lambda\sigma^2 + e_0^2)/(\lambda(1 + c) - c)$.

Proof. Appendix A. Elementary but complete: concavity gives a unique downcrossing of $T(v) - v$; contraction above the fixed point plus a monotone barrier below give convergence under vanishing perturbations. $\square$

Remark 1 (why annealing cannot help singular data). Truncation enters (2) only through the additive, schedule-killable s_g^2 . The plateau v_∞ is determined by e^2 alone, which no schedule touches. The remedy of Khelifa et al. [1] is the special case $e^2 \equiv 0$ (their Case 1 assumptions make the estimation term schedule-dominated), where indeed $v_\infty = \sigma^2$. Case 2 is precisely the regime where e^2 has a floor, the subject of the next section.

5 The capacity floor

We now prove that for manifold data the sampling excess e^2 of any truncated reverse-SDE sampler is floored by the normal slope its ε -predictor realises. No decomposition of the network is required. Let X_t be the sampler's *own* ensemble at noise level t and define the *realised normal slope*

$$\hat{\kappa}(t) := \frac{\text{Cov}(X_t, \hat{\varepsilon}(X_t, t))}{\text{Var}(X_t)}, \quad (3)$$

the L^2 regression slope of the estimator on the normal coordinate under the law the sampler actually visits. For an affine estimator this is its slope; for an arbitrary nonlinear one it is the only functional of $\hat{\varepsilon}$ that the variance flow can see. The structural fact behind the floor (proof in Appendix B):

Lemma 1 (variance ODE; we prove, exact). *Under the reverse Euler scheme, for an arbitrary estimator $\hat{\varepsilon}$ and with $\hat{\kappa}$ as in (3), the ensemble normal variance obeys, in the small-step limit,*

$$-\frac{dV}{dt} = \left(1 - \frac{2\hat{\kappa}(t)}{s(t)}\right)V + 1. \quad (4)$$

With the exact slope $\hat{\kappa} = \kappa$, $V(t) = \gamma^2(t)$ solves (4) exactly. The forcing +1 is the sampler’s own injected noise. No linearisation is performed: to first order in the step a single Euler step moves the variance only through $\text{Cov}(X, \hat{\varepsilon}(X))$, the $O(dt^2)$ remainder vanishing in the limit; equivalently, in continuous time Itô’s formula yields (4) directly, with no expansion. The identity is exact for a normal coordinate that is itself a plain Itô diffusion, which a flat normal coordinate is. A curved one is not, and the correction has a closed form. Write h for the signed distance to the manifold, so $|\nabla h| = 1$ and $n := \nabla h$ is the unit normal; Itô gives $dh = \langle n, f \rangle d\tau + \frac{1}{2}\Delta h d\tau + \langle n, dW \rangle$, and the same two lines yield, still exactly and still for an arbitrary $\hat{\varepsilon}$,

$$-\frac{dV}{dt} = \left(1 - \frac{2\hat{\kappa}(t)}{s(t)}\right)V + 1 + \text{Cov}(h, \Delta h) + [\text{Cov}(h, \langle n, x \rangle) - V]. \quad (5)$$

The bracket vanishes for a sphere about the origin and for an affine subspace through it, which between them cover every geometry we analyse: the segment and the eight out-of-plane axes of the codimension-9 ring are affine subspaces through the origin, and the ring tube and S^2 are spheres about it. The operative term is $\text{Cov}(h, \Delta h)$, and $\Delta h \equiv 0$ for a flat normal coordinate — which is why (4) is exact there, rather than by assumption. For a sphere of radius R in $\mathbb{R}^D$, $\Delta h = (D-1)/r$ and

$$\text{Cov}(h, \Delta h) = (D-1) \text{Cov}(r, \frac{1}{r}) = (D-1)(1 - \mathbb{E}[r] \mathbb{E}[1/r]) = -\frac{D-1}{2} \mathbb{E}\left[\frac{(r-r')^2}{r r'}\right] \quad (6)$$

for an i.i.d. copy r' : negative by Cauchy–Schwarz, and unbounded relative to its own linearisation $-(D-1)V/R^2$ as soon as r carries mass at small radius. The linearisation is already $2.09\times$ wrong for a plain Rayleigh law in $D=2$. We verify (5) against a case with a closed form: for Brownian motion from the origin in $\mathbb{R}^D$, $r = \sqrt{\tau} \chi_D$ gives $dV/d\tau = D - c_D^2$ with $c_D = \sqrt{2} \Gamma(\frac{D+1}{2})/\Gamma(\frac{D}{2})$, and the difference between that and (5) reduces to $z\Gamma(z) = \Gamma(z+1)$ at $z = (D-1)/2$ — identically zero, confirmed to 50 digits for integer and non-integer D alike, and reproduced to 100.6% by direct simulation under the sampler’s own Euler scheme (`radial_ode_identity.py`).

The same Itô step fixes the first moment, where the curvature term is far larger and does not vanish as the ensemble concentrates:

$$-\frac{d\mu}{dt} = \frac{\mu}{2} - \frac{\mathbb{E}[\varepsilon_n]}{s} + \frac{1}{2}\mathbb{E}[\Delta h], \quad \mathbb{E}[\Delta h] = (D-1)\mathbb{E}[1/r] \text{ on a sphere, } 0 \text{ on a flat coordinate.} \quad (7)$$

Being a different moment of the same correction, (7) is an independent check rather than a restatement, and it is the sharper one: $\text{Cov}(r, 1/r) \rightarrow 0$ as the tube contracts whereas $\mathbb{E}[1/r] \rightarrow 1/(\alpha R)$ stays $O(1)$.

We also confirm (5) on the trained ring network itself, which requires abandoning the estimator we used previously. Differentiating $V(t)$ along a sampler trajectory forms $(V_{i+1} - V_i)/dt$, so sampling error in V is divided by dt and diverges as the grid refines: within one script the residual so measured moved from -0.041 at 1000 steps to -0.050 at 2000 (the earlier probe, at 4000 steps and a different

sample count, reported -0.062), and the *predicted* term moved 10.6% over the same refinement, because the trajectory itself changes with the step count. Instead we freeze the ensemble at a given t and take a single Euler step, paired and antithetic so that the $O(\sqrt{d\tau})$ sampling term cancels exactly (`curved_ode_onestep.py`). The flat and curved forms then differ by a measurable margin and the measurement selects between them. At $t = 0.5$ the flat right-hand side is -0.316 against a measured -0.495 ; (5) predicts -0.496 . Sweeping t at a fixed step $d\tau = 10^{-4}$, the raw ratio is 0.993, 0.992 and 0.978 at $t = 0.5, 0.2$ and 0.1 , and then degrades as the term shrinks toward the estimator's own error.

That error is measurable rather than guessed, because the flat segment runs the identical estimator on a geometry where (6) vanishes *identically*: whatever residual it reports is the method's bias at that t and step. It reports $-0.0003, +0.0011, +0.0013, +0.0019, +0.0031, +0.0051$ at $t = 0.5$ down to 10^{-2} , growing as $1/s$ — the Euler drift bias, which the two geometries share. Subtracting it, the ring's ratio is 0.992, 1.010, 1.021, 0.994, 0.980 and 1.078 across the same range: the identity holds to within 2% at five of those six times and 8% at the sixth, over a $50\times$ span in t , with at most 0.0013 left unexplained anywhere (`onestep_control_corrected.py`). The correction is a genuine control and not a fitted offset: at $d\tau = 10^{-5}$, where the $O(d\tau)$ bias is gone and only random noise remains, subtracting the flat geometry's independent draws stops helping, exactly as it should.

$S^2 \subset \mathbb{R}^3$ then tests the coefficient itself, at $(D - 1) = 2$ rather than 1, on a different network, a different radial law and a different ambient dimension (`curved_ode_onestep_sphere.py`). The variance equation gives ratios 1.000, 1.002, 0.996, 0.967 at $t = 0.5$ to 0.05. The mean equation (7) is the decisive one, because its curvature term does not shrink: across all eight probe times from $t = 0.5$ down to 2×10^{-3} the ratio is 0.9994, 0.9986, 0.9997, 1.0003, 1.0011, 1.0003, 1.0032, 1.0068 — every one within 0.7%, over a $250\times$ span in t . Over the same range, forcing the coefficient to 1 instead of $D - 1 = 2$ misses by 1.999 to 2.014, so the scaling is confirmed and not fitted. The margin being tested is not marginal: at $t = 0.05$ the flat form predicts $+0.099$ against a measured $+0.505$, and at $t = 5 \times 10^{-3}$ it predicts -0.711 against a measured -0.312 . On a curved normal coordinate the flat mean equation is wrong by factors, and (7) is right to a fraction of a percent.

Because Δh enters only through a covariance with h , Remark 2 converts (5) into a bound with no further assumptions.

Corollary 1 (the floor survives curvature). *If Δh is L -Lipschitz in h on the support of the ensemble, then $|\text{Cov}(h, \Delta h)| \leq LV$ and*

$$-\frac{dV}{dt} \geq \left(1 - \frac{2\bar{\kappa}}{s} - L\right)V + 1,$$

so Theorem 2 holds verbatim with the coefficient shifted by $-L$. The hypothesis is a statement about the medial axis: Δh measures the mean curvature of the level sets of the distance function, and it is exactly on the medial axis, where the nearest-point projection stops being single-valued, that it fails to be bounded. The constant is then supplied by the reach. For M of dimension d and codimension one with reach τ , the level sets of h have shape operator $S(I - hS)^{-1}$ with $\|S\| \leq 1/\tau$, so

$$|\Delta h| \leq \frac{d}{\tau - |h|} \quad \text{on the tube } |h| < \tau, \quad \text{hence} \quad L \leq \frac{d}{(\tau - h_{\max})^2}, \quad (8)$$

with $h_{\max}$ the furthest the ensemble strays from M .

(8) is an equality for a sphere, where $\tau = R$, $d = D - 1$ and $|h| = |r - R|$ return $\Delta h = (D - 1)/r$ exactly, so nothing is lost by stating it this way; for the ring ($\tau = 2.5$, $d = 1$, $h_{\max} = 1.0$) it gives $L = 0.444$, the same constant as before. What is gained is that τ is a property of the manifold rather than of the sampler's ensemble, and it is the quantity this literature already assumes. We

check (8) on a geometry whose reach is not a radius (`reach_bound_check.py`): for an ellipse with semi-axes 3 and 2, whose reach is $B^2/A = 1.33$ and whose medial axis is a segment rather than a point, the bound holds at every test point inside the tube with worst ratio 1.000, attained at the end of the major axis where the curvature is maximal. Past the reach a finite-difference Laplacian of h diverges to $\sim 2 \times 10^4$, which is the medial axis making itself visible: there h is not C^2 and no bound of this shape is claimed.

The shift is a *constant* while $2\bar{\kappa}/s$ diverges as $s \rightarrow 0$, so curvature is irrelevant exactly where the floor is set: for $D = 2$ and $a = 1.5$ it is 4.5% of the coefficient at $t = 1$, 0.57% at $t = 10^{-2}$ and 0.018% at $t = 10^{-5}$. Near the floor ($t \in [10^{-4}, 10^{-2}]$) the ring’s sampler ensemble has median $\mathbb{E}[r] = 2.46$ with its 0.1% radial quantile at 1.32–1.45, giving $L \approx 0.5$ and lowering $\Phi_{\text{det}}(3.9)$ by 2.0–2.6% (`curved_floor_envelope.py`, which reproduces `deficit_floor_law.integrate` exactly at $L = 0$). This is what reconciles a large measured ODE residual on the ring with a floor bound the ring nonetheless satisfies: the residual lives at moderate t and the floor does not.

We flag the one real weakness of this route rather than leave it to be found. L is set by an essential infimum, so it is governed by rare small-radius stragglers and not by the bulk of the ensemble. Reading a off the 0.01% quantile instead (0.36–0.46 near the floor) costs 12–19%, and an ensemble genuinely reaching $a = 0.05$ would push the bound below σ^2 and assert nothing at all. The exact term (6) has no such sensitivity — it is two measured moments, $\mathbb{E}[r]$ and $\mathbb{E}[1/r]$ — so where a numerical floor is wanted on curved data we would integrate with the measured $\mathbb{E}[r]\mathbb{E}[1/r]$ and use Corollary 1 only for the assumption-light statement. We note that the same geometric quantity carries the assumptions elsewhere in this literature: Azangulov et al. [19] obtain ambient-dimension-free rates for DDPMs on a manifold whose reach is bounded below, the reach being precisely the distance from the manifold to its medial axis. A distance-from-the-medial-axis hypothesis of the shape above is therefore the natural one to state, and it is a strictly weaker requirement than curvature bounds, since a highly curved manifold can have large reach and a nearly flat one can have small reach where two sheets approach.

We should be precise about how far that parallel goes, because the two uses of the reach are not the same. Azangulov (personal communication) points out that their score-regularity result is in fact reach-agnostic: what the proof uses is that the support decomposes into charts $f_i(B_d(0, 1))$ with each f_i Lipschitz, and a lower bound on the reach enters only as a convenient way to guarantee such a decomposition exists. Our use is not substitutable in that way. (8) bounds the shape operator of the level sets of h , and it is the reach itself, not a chart decomposition, that supplies $\|S\| \leq 1/\tau$. So the reach is sufficient in both settings and necessary in neither obvious sense, but it is load-bearing here in a way it is not there, and we would expect our constant to survive a weakening of the hypothesis only if the weaker condition still controls Δh .

Remark 2 (a checkable sufficient condition). The hypothesis of Theorem 2 constrains $\hat{\kappa}$, a regression slope. A pointwise bound implies it: if $\hat{\varepsilon}(\cdot, t)$ is $\bar{\kappa}$ -Lipschitz in the normal coordinate then, for an i.i.d. copy X' of X_t ,

$$\text{Cov}(X, \hat{\varepsilon}(X)) = \frac{1}{2}\mathbb{E}[(X - X')(\hat{\varepsilon}(X) - \hat{\varepsilon}(X'))] \leq \frac{\bar{\kappa}}{2}\mathbb{E}[(X - X')^2] = \bar{\kappa}\text{Var}(X),$$

so $\hat{\kappa}(t) \leq \bar{\kappa}$ under *any* law. Lipschitz control is therefore sufficient but not necessary; we state the theorem on $\hat{\kappa}$ because it is the weaker hypothesis and because it is what (4) actually contains.

This is the precise sense in which the floor is *not* an artifact of linearizing the score. An earlier version of this argument split $\hat{\varepsilon}$ into an affine $L^2(p_t)$ fit plus an orthogonal residual and invoked the normal equations to discard the residual to first order. That step is no longer needed, and with it goes its one soft spot: the normal equations hold under p_t , whereas the sampler’s own ensemble

deviates from p_t . Equation (4) is exact for the fully nonlinear network precisely because $\hat{\kappa}$ is defined under the sampler's law from the outset. The crux is the geometry of the required slope:

Lemma 2 (finite band; we prove). $\kappa(t)$ rises from 0, peaks at $\kappa_{\max} = 1/(2\sigma)$ at $t \approx \sigma^2$, and falls back to 0 as $t \rightarrow 0$. Hence for any ceiling $\bar{\kappa} < \kappa_{\max}$ the deficit set $\{t : \kappa(t) > \bar{\kappa}\}$ is a finite band; below it the exact slope is representable again, but only $2 \log(1 + \rho^2/4) \leq \rho^2/2$ contraction e -folds remain ($\rho := 2\sigma\bar{\kappa}$), so variance frozen in the band persists to $t_0 \rightarrow 0$.

In the singular limit the whole flow collapses onto a one-parameter family (Appendix C for proofs and error terms):

Lemma 3 (singular-limit ODE; we prove). Rescale $t = \sigma^2 u$, $V = \sigma^2 W$, and write the slope envelope as $\hat{\kappa}(t) = (1/\sigma)m(u)$ with $m(u) = \rho/2$ (unconditional envelope) or $m(u) = \min\{\sqrt{u}/(1+u), \rho/2\}$ (no-overshoot envelope). Then (4) becomes, as $\sigma \rightarrow 0$ at fixed ρ ,

$$-\frac{dW}{du} = 1 - \frac{2m(u)}{\sqrt{u}} W, \quad (9)$$

a σ -free equation; the uncapped equation has exact solution $W = 1 + u$ (the rescaled true marginal), and deviations contract quadratically under downward integration. Consequently $\lim_{t_0 \rightarrow 0} V(t_0)/\sigma^2 = g_{\text{class}}(\rho)(1 + O(\sigma^2/\rho^2))$ (a relative correction: the boundary-layer perturbation scales with the leading term $W \sim \rho^{-2}$, so the additive gap to g is $O(\sigma^2/\rho^4)$, as the finite- σ values confirm): **universality in ρ is a theorem.**

How to read the theorem. There is one dimensionless knob, $\rho = \bar{\kappa}/\kappa_{\max} = 2\sigma\bar{\kappa}$: how close the network's slope ceiling gets to the peak slope the data demands. The floor is a universal function of ρ alone. Version (U) makes no assumption on the network and gives the clean bound $1/(8\bar{\kappa}^2)$; version (N) additionally assumes the network never over-shoots the true slope (which is what we measure) and gives a slightly tighter constant. The two constants differ by only 5.5%, so for intuition either one is “the” floor. The single takeaway: when $\rho < 1/\sqrt{2}$ the floor sits strictly above the true width σ^2 , i.e. the sampler is stuck too fat, and it gets relatively worse the thinner the manifold.

Theorem 2 (closed-form capacity floor; we prove). Let the trained score's normal slope satisfy $\hat{\kappa}(t) \leq \bar{\kappa} < \infty$ (a finite normal Lipschitz constant: a bound on slope, not on parameter count), and set $\rho = 2\sigma\bar{\kappa}$. Then for every truncation time $t_0 > 0$ and every schedule:

(U) **Unconditional.** With no further assumption (arbitrary overshoot permitted),

$$V(t_0) \geq \sigma^2 \frac{1}{2\rho^2} (1 + o(1)) = \frac{1}{8\bar{\kappa}^2} (1 + o(1)) \quad (\sigma \rightarrow 0), \quad (10)$$

where $1/(2\rho^2)$ is the exact value of (9) with $m \equiv \rho/2$: $W(0^+) = \int_0^\infty e^{-2\rho\sqrt{u}} du = 1/(2\rho^2)$. The floor exceeds σ^2 whenever $\rho < 1/\sqrt{2}$, i.e. $\bar{\kappa} < \kappa_{\max}/\sqrt{2}$.

(N) **No-overshoot.** If additionally $\hat{\kappa}(t) \leq \kappa(t)$ (the empirically observed profile; the network never over-contracts), the tighter floor $V(t_0) \geq \sigma^2 g(\rho)(1 + o(1))$ holds with the explicit closed form: writing $q = \sqrt{1 - \rho^2}$, $x_\mp = (1 \mp q)/\rho$, $u_\mp = x_\mp^2$,

$$W_{\text{exit}} = (1 + u_+)e^{-4q} + \frac{(3 - 2q) - (3 + 2q)e^{-4q}}{2\rho^2}, \quad g(\rho) = 1 + \frac{W_{\text{exit}} - (1 + u_-)}{(1 + u_-)^2}. \quad (11)$$

$g(\rho) > 1$ for every $\rho \in (0, 1)$ with $g \rightarrow 1$ only as $\rho \rightarrow 1^-$, and $g(\rho) = C_0/\rho^2 + O(1)$ as $\rho \rightarrow 0$ with $C_0 = (1 + 3e^{-4})/2 = 0.5275$, i.e. the deterministic single-pass floor $\Phi_{\text{det}} := \lim_{t_0 \rightarrow 0} V(t_0) = \sigma^2 g(\rho)$ satisfies $\Phi_{\text{det}} \rightarrow ((1 + 3e^{-4})/8)\bar{\kappa}^{-2} = 0.1319 \bar{\kappa}^{-2}$.




Figure 2: **The two-horned trap (*we observe*)**. Self-consuming recursions at fixed $t_0 = 10^{-4}$, no anchor, 3 seeds. With the reverse-SDE noise on, the recursion sticks fat at 10.8 to $13.2 \sigma^2$ (red). With the noise zeroed (deterministic sampler), it collapses to a *single point* within one generation (orange; verified loss of all 2-D spread, no numerical overflow). Neither branch reaches the tube (σ^2 , dotted); the anchored recursion (blue, Section 6) holds it.

The two asymptotic constants, $1/8$ and $(1 + 3e^{-4})/8$, differ by 5.5%: the class distinction is quantitatively minor. The closed forms match direct integration of (9) to $\leq 2.4 \times 10^{-5}$ across $\rho \in [0.1, 1]$, and the finite- σ ODE converges to them from above at the stated relative $O(\sigma^2/\rho^2)$ rate.

Corollary 2 ($\sigma \rightarrow 0$: the answer to Case 2; *we prove*). *Hold $\bar{\kappa}$ fixed and let $\sigma \rightarrow 0$ ($J \rightarrow \infty$). Then unconditionally $V/\sigma^2 \geq (1 + o(1))/(8\bar{\kappa}^2\sigma^2) \rightarrow \infty$: a truncated sampler held at a fixed slope ceiling collapses relatively without bound, for every annealing schedule. Conversely, bounded relative collapse $V/\sigma^2 \leq C$ requires*

$$\bar{\kappa} \geq \frac{1}{2\sigma\sqrt{2C}} = \Omega(1/\sigma) = \Omega(\sqrt{J}). \quad (12)$$

Truncated sampling of singular data demands a slope budget growing with the singularity. This constrains the normal Lipschitz norm: a fixed architecture can meet it only by enlarging its weights, and whether training in fact does so is the empirical question the $\bar{\kappa}(w)$ probe addresses. We thank Wei Huang for pressing this distinction.

What each ingredient does (*we observe*). The ablation of Figure 2 separates the roles exactly as (4) dictates: the sampler’s own noise (the +1) is load-bearing for the floor’s *existence*: switching it off sends $V \rightarrow 0$ (point collapse, the deterministic horn), while the finite ceiling $\bar{\kappa}$ makes the floor *larger than σ^2* (the stochastic horn). The two branches bracket the target from both sides; no choice of the sampler noise magnitude attains it.

Confronting the floor with a real network (*we observe*). For our standard pipeline the measured slope profile tracks $\kappa(t)$ faithfully at benign times ($\hat{\kappa}/\kappa = 0.95$ to 0.97 for $t \geq 0.1$) and then

saturates while κ continues to ~ 10 . Which ceiling one reads off depends on the law the regression is taken under, and (3) settles that: the hypothesis is the slope under the sampler’s *own* ensemble. Probing the forward process instead returns a systematically larger ceiling. Measured on the same networks across five tube widths (`kbar_sampler_law.py`; $\sigma = 0.03$ to 0.13), the sampler-law ceiling is 0.83 to 0.87 of the forward-process one, monotone in σ and without exception.

What the ceiling actually depends on (we observe). We had first attributed the gap to the sampler’s ensemble being fatter than p_t and so probing the response into its flattening tails, and then, on seeing the flat geometries agree, rejected that explanation in favour of curvature. Both readings were too simple. Sweeping the ensemble width at fixed network and fixed t (`slope_vs_width.py`) shows the measured ceiling falling with width on *every* geometry — by 10 – 28% over a sixfold sweep on the flat segment and 23 – 55% on the ring — so saturation is real and curvature roughly doubles the sensitivity to it rather than supplying the effect on its own. What then distinguishes the geometries is not their response to width but how far their sampler ensembles actually depart from p_t , and in what way. Measured on the sampler’s own trajectory (`ensemble_shape.jsonl`), the flat segment stays within 1.06 – $1.43\times$ the width of p_t and symmetric (skewness $+0.05$ to -0.03), which predicts a probe ratio of 0.996 against 0.995 – 0.998 measured: the flat case is quantitatively closed. The ring reaches $2.05\times$ the width, sits 0.03 – 0.04 *inside* the contracted manifold, and is strongly left-skewed, from -0.44 at $t = 0.1$ to -1.02 at $t = 5 \times 10^{-3}$. Skewness is the sharpest *descriptive* discriminator: essentially zero on flat normal coordinates and order one on curved ones. It is not, however, the operative one; see (14). Injecting a comparable skew into an otherwise matched synthetic ensemble moves the measured ceiling from 3.47 to 3.37 . Width and skew supply 5.6% of the ring’s $3.6 \rightarrow 3.0$ gap; moments beyond the third and the coupling between radius and angle supply the rest, and the budget closes. The methodological point is the one we would ask a reader to take: $\bar{\kappa}$ is not a property of a trained network alone. It is a property of a network *and* the ensemble it is read on, and any reported score-slope or Lipschitz ceiling that does not name its measurement ensemble is under-specified. The geometries of §7 make the point sharply. Re-measuring all of them under both probes (`capacity_domains_samplerlaw.py`) gives a probe ratio of 0.995 – 0.998 on the straight segment and 0.975 – 0.988 on the codimension-9 ring, whose slope is read on its eight *out-of-plane* axes: both flat normal coordinates, both agreeing to within 1% . The sampler ensemble is fatter in those settings too, so fattening alone cannot be the mechanism. Where the normal coordinate is *curved* the probes separate: 0.830 – 0.841 on the ring tube, whose radial normal curves in one direction, and 0.740 – 0.746 on S^2 , which curves in two. The ordering is consistent across all five widths, though the sphere’s deficit is 1.5 – $1.8\times$ the ring’s rather than the $2\times$ strict proportionality in the number of curved directions would give, so we report the ordering and not a law. Two consequences. The flat tube the theorem idealises is exactly the case where the two readings coincide, so (3) and the forward-process probe agree in the setting Theorem 2 assumes; and the discrepancy on ring data is a curvature artifact of the probe, of the same origin as the $O(s^2/\tau)$ term §8 declines to carry. Reassuringly, the relative degradation is untouched: the drop across $\sqrt{w}$: $0.03 \rightarrow 0.13$ reads 42.3% against 42.2% on the segment, 29.8 against 29.2 on the sphere, and 48.7 against 48.0 on the codimension-9 ring, so every trend claim in §7 stands under either probe.

That last point is not a coincidence of these three geometries but an identity. Writing r for the probe ratio and $D = 1 - \bar{\kappa}(w_{\max})/\bar{\kappa}(w_{\min})$ for the relative drop,

$$D_{\text{samp}} = 1 - \frac{r(w_{\max})}{r(w_{\min})}(1 - D_{p_t}), \quad (13)$$

so the drop is invariant to the probe up to the *variation* of r across the width range, and is wholly insensitive to its level. (13) reproduces all four measured drop shifts to 0.005 percentage points.

Across the four geometries measured on a common protocol — flat, singly curved and doubly curved — r varies by at most 1.33% within a geometry, so a degradation law measured under p_t transfers to (3) with a shift of under one point. The MNIST 8×8 degradation of 21.1% would then read 20.0–21.0%, and we do not re-run it. This bound is conditional and we state the condition: the $t_0 = 10^{-4}$ ring probe of §5, measured at one seed per width, varies by 4.97% instead, which would carry the same drop to 17.2%. We take the four-geometry figure as the better-sampled estimate and flag the single-seed one as the reason the transfer is an estimate rather than a measurement. We flag the assumption rather than hide it: on real data the normal frame is a per-point kNN-PCA estimate, and a sampler-law probe would need those frames anchored to off-manifold points, a choice the theory does not force and which we prefer not to make silently. Since the floor scales as $\bar{\kappa}^{-2}$, the quantity the theorem constrains gives the *larger* bound. Over six seeds at $\sigma = 0.05$: $\bar{\kappa} = 3.26$ rather than 3.90, so the envelope Φ_{det} reads $5.49\sigma^2$ rather than $3.83\sigma^2$ (Figure 3), accounting for 81% of the measured single-pass floor against 58% read under p_t , with no seed falsified. The correct reading is also the more reproducible one: its seed spread is 77–85% against 42–72% for the forward-process probe, a factor 4.2 narrower, which is what one expects once the probe and the dynamics are evaluated under the same law. The exact-score control reproduces $\gamma^2(t_0)$ to all displayed digits.

The finite- σ error term (we resolve). Theorem 2 is a $\sigma \rightarrow 0$ statement carrying a *relative* correction $O(\sigma^2/\rho^2)$. Comparing its envelope against finite- σ data without applying that correction is an error, and it is the one we made: at $\sigma = 0.13$ the envelope exceeds the measured floor in three seeds of six, by at most 3.4%, while the theorem’s own error term at that width ($\rho = 0.638$) permits 4.15%. Every width, worst seed included, is consistent once the stated correction is applied: 2.2% allowed against 0% observed at $\sigma = 0.03$, rising to 4.2% against 3.4% at $\sigma = 0.13$. There is no violation and no restricted validity range; the bound tightens toward equality as the tube fattens because ρ grows, which is what the theorem predicts.

A second and independent reason the comparison was never in tension is that Φ_{det} , integrated from (4), is not the correct bound on a ring at all. These networks are trained on a curved manifold, so the governing equation is (5) and the bound is the one shifted by $-L$ (Corollary 1). That shift lowers the envelope by 2.4% at $\sigma = 0.03$ rising to 4.3% at $\sigma = 0.13$, and it removes the excess outright rather than tolerating it: at $\sigma = 0.13$ the worst seed falls from 103.4% of its measured floor to 98.9% taking $a = 1.5$, and to 97.8% at the ensemble’s measured 0.1% radial quantile $a = 1.32$. The conclusion does not turn on that choice — every a below 1.75 suffices, and the sampler’s own ensemble puts 0.1% of its mass at radius 1.32 and 0.01% at 0.363. Under the corrected bound no seed at any width is exceeded, and we report the two resolutions separately because they are independent: the first is a tolerance the theorem already carries, the second is a statement that the object being compared was the wrong one. We record the curvature term we checked along the way and rejected as too small. The radial Itô correction $(D - 1)/(2r)$ adds a contraction to (4) on curved data; we first estimated it by linearising $1/r$ about the mean radius, as $-(D - 1)V/R^2$, and rejected it at 0.7–1.7% of the envelope. The linearisation was wrong — the exact term is (6), which is $2.09\times$ larger already for a Rayleigh law and unbounded relative to the linearisation when r carries mass at small radius — but the conclusion was not. Corollary 1 puts the rigorous cost at 2.0% of $\Phi_{\text{det}}(3.9)$ for an ensemble staying at radius ≥ 1.5 , still an order short of the discrepancy, and the shift is a constant against a $2\bar{\kappa}/s$ that diverges, so it is 0.018% of the coefficient by $t = 10^{-5}$. Curvature does not set the floor; the moment structure of the ensemble does.

The probe gap, in closed form (we prove, and measure). The four channels are not independent effects but terms of a single expansion. For any ensemble q and any fixed estimator, expanding $\hat{\epsilon}$

about the ensemble mean m and taking the regression slope gives

$$\hat{\kappa} = \hat{\varepsilon}'(m) + \frac{1}{2}\hat{\varepsilon}''(m)\frac{\mu_3}{\mu_2} + \frac{1}{6}\hat{\varepsilon}'''(m)\frac{\mu_4}{\mu_2} + \dots, \quad (14)$$

with μ_k the central moments of the normal coordinate under q . An affine estimator kills every term past the first and $\hat{\kappa}$ becomes ensemble-free; a nonlinear one does not, and the two probes then differ because they carry different m , μ_3 and μ_4/μ_2 . Evaluating (14) on the trained networks, with the derivatives taken by finite differences of the angle-averaged radial response and the moments read off the sampler’s own ensemble, reproduces the measured ceiling to 1.3% on the flat segment and 1.8% on the ring at $t = 5 \times 10^{-3}$.

The expansion also identifies which term matters, and it is not the one we first supposed. The skewness contribution is negligible on both geometries: it moves the ring’s ceiling from 3.598 to 3.597. The work is done by the *fourth*-moment term weighted by the network’s third derivative, which moves the segment by 1% ($4.777 \rightarrow 4.733$) and the ring by 18% ($3.598 \rightarrow 2.945$). Curvature enters through $\hat{\varepsilon}'''$, the sharpness with which the learned radial response saturates; the skew we measured is a correlate of that sharpness rather than its cause. This is what makes $\bar{\kappa}$ ensemble-dependent in a way no low-order summary repairs: the leading correction is a fourth moment against a third derivative, and neither is recoverable from a reported width.

The same budget, channel by channel (we observe). The margin narrows monotonically as the tube fattens: 77–85% of the measured floor at $\sigma = 0.05$, 90–97% at $\sigma = 0.10$, and 97–103% at $\sigma = 0.13$, all inside the relative error term above, and 75–82%, 87–93% and 93–99% under the curved-geometry bound of Corollary 1, which no seed exceeds.

The remaining gap between envelope and measurement decomposes completely. Taking the ring at $t = 5 \times 10^{-3}$ and moving one feature at a time from the p_t ensemble to the sampler’s own: widening to the sampler’s width costs 2.6% of the slope; matching its skew a further 3.0%; replacing the moment-matched surrogate by the sampler’s actual radial marginal, so that moments beyond the third enter, a further 7.2%; and restoring the coupling between radius and angle, which an isotropic surrogate destroys, a final 7.3%. The product $3.564 \times 0.974 \times 0.970 \times 0.928 \times 0.927 = 2.90$ against 2.894 measured, so the four channels close the budget rather than half of it (`gap_budget.jsonl`). Two of them are invisible to any summary statistic of the ensemble: the sampler’s radial law is not determined by its first three moments, and it is not a product of a radial marginal with a uniform angle. That is the concrete sense in which $\bar{\kappa}$ is a property of the ensemble and not of the network. It is not evidence against Theorem 2, whose chain we verify symbolically (Appendix C); it is the flat-tube idealisation expiring where Section 8 says it will, curvature entering at $O(s^2/\tau)$ and not carried in the constants. At $\sigma = 0.13$ on a radius-2.5 ring that correction is no longer negligible. The monotone narrowing is the signature of an asymptotic bound pushed past its regime, not of a defect in the derivation, and it places the honest validity edge of the *numerical* comparison between $\sigma = 0.10$ and $\sigma = 0.13$. We checked the obvious alternative first: taking $\bar{\kappa}$ as the maximum over the whole deficit band rather than over five probe times moves it by <0.02 , and leaves two of six seeds violating.

What is left over (*we observe*). Re-running the sampler’s own discrete variance recursion with measured inputs now resolves the remainder into *one* channel rather than two. Earlier versions of this measurement carried an **ensemble-displacement** term, worth $+2.6\sigma^2$, for precisely the probe mismatch just described; under (3) that is not a correction to the bound but part of the hypothesis, and it is no longer counted separately. What remains is **profile shape**: the measured $\hat{\kappa}(t)$ sags ≈ 0.5 – 0.6 below its own ceiling inside the deficit band, worth $+1.1$ to $+1.4\sigma^2$ at the data width. Both figures are read under (3) for consistency with the hypothesis; under p_t the same two

quantities are 0.6–0.7 and +0.5 to +0.9 σ^2 , and the channel is worth roughly half as much because the p_t ceiling sits higher and the envelope correspondingly lower. Read the consistent way the single channel carries the whole remainder: the sampler-law envelope is 5.49 σ^2 against a measured floor of 6.80 σ^2 , a gap of 1.31 σ^2 , and profile sag alone supplies 1.1–1.4 of it. (The two runs use different seed sets, so this is a magnitude check, not an exact reconciliation.) Beyond it the affine recursion driven by the measured ensemble slope reproduces the sampler to $-0.4 \sigma^2$: the nonlinear remainder is small. **Interventional test.** The law survives moving the ceiling: five training protocols (t -importance-sampling into the band, $4\times$ steps, a higher-frequency time embedding, and combinations) move $\bar{\kappa}$ from 3.2 to 5.8 read under (3) (3.8 to 7.6 under p_t); the measured floor drops from 6.4–7.2 σ^2 at the two low- $\bar{\kappa}$ protocols into the 3.0–3.4 σ^2 range at the three high- $\bar{\kappa}$ ones (a group-level fall with protocol-dependent scatter, not a clean $1/\bar{\kappa}^2$ trace), and — the load-bearing statement — every one of the five sits strictly above its own proven bound.

That statement is the one most exposed by (3), since a lower sampler-law ceiling *raises* each arm’s bound, so we re-measured all five under both probes (`ceiling.origin.samplerlaw.py`, three seeds per arm). It survives: no arm is violated, and the worst now sits at 78% of its measured floor against 60% read under p_t . The probe ratio is not a constant of the geometry but a property of the trained network, running 0.754 to 0.853 across the five and splitting cleanly by training protocol — 0.843–0.853 for the three arms drawing noise levels uniformly, 0.754–0.769 for the two importance-sampling half their draws into the deficit band, which trains the network harder exactly where the sampler’s ensemble is fattest relative to p_t . A single scale factor is therefore not a substitute for the measurement.

What sets the ceiling (*we observe*). The interventions also identify the ceiling’s origin. Raising the time-embedding frequency alone does nothing ($\bar{\kappa} = 3.5$): expressivity in t is not binding. Importance-sampling training times into the band ($\bar{\kappa} \rightarrow 6.7$) or simply training $4\times$ longer ($\bar{\kappa} \rightarrow 6.4$) moves it strongly: under the uniform t -draw the deficit band holds $<1\%$ of the training measure (and $t < 10^{-3}$ is never drawn at all, though the sampler visits it), so the ceiling is *gradient starvation* — an optimization constant, not an architectural one, consistent with its independence of sample size and its degradation on fatter pools (a fatter tube spreads the same starved gradient over an easier, wider band). The Case-2 conclusion is unchanged — as $\sigma \rightarrow 0$ the demanded peak $1/(2\sigma)$ diverges, so any fixed training budget caps $\bar{\kappa}$ — but the constant is steerable: band-weighted noise-level sampling buys a 2–3 $\times$ lower per-pass floor at fixed size. It cannot reach σ^2 , and the recursion compounds whatever remains, so anchoring is still the fix.

An honest gap, closed in three measured steps (*we model*). Iterating the floor through the recursion (every generation’s pool is at least as wide as $\lambda\sigma^2 + (1 - \lambda)\Phi_{\text{det}}$) with the network’s *clean-tube* ceiling $\bar{\kappa} = 3.9$ gives a deterministic fixed point $v_{\text{det}}^* = 4.5 \sigma^2$. The measured recursion floor is $\approx 11 \sigma^2$, a factor 2.4 above; our theorem claims only a *lower bound* on the recursion (respected). We account for that factor in three measured steps, each a directly logged measurement.

Step 1, capacity degradation. The single-width probe (Appendix F) shows the ceiling itself degrades on fatter training pools, from 3.2–3.6 at $\sqrt{w} = 0.03$ down to 2.4–2.6 at $\sqrt{w} = 0.13$ read under (3) (4.0–4.4 down to 2.8–3.1 under p_t ; ranges over three training seeds, `poolwidth.probe.samplerlaw.jsonl`), and the degradation replicates under a half-width network, under a different ring radius (slopes -9.4 to -13.0 across all five measurements), across three further synthetic geometries — a flat segment (the flat-tube idealization made exact, so the degradation does not require curvature; 42% drop over the same width range), a sphere in $\mathbb{R}^3$ (30%), and a circle in $\mathbb{R}^{10}$ (codimension 9; 49%) — and, crucially, on *genuine MNIST* 8×8 pixel data (whose digit manifold has a measured intrinsic




Figure 3: **The capacity floor: predicted, measured, and moved.** Blue line with markers: finite- σ variance-ODE envelope floor at $\sigma = 0.05$ (no fit). Dashed: the exact $\sigma \rightarrow 0$ law $\frac{(1+3e^{-4})}{8}\bar{\kappa}^{-2}$ of Theorem 2. Stars: the envelope bound at the measured ceiling, read two ways. Hollow, under the forward process ($\bar{\kappa} = 3.9$, bound $3.8\sigma^2$); filled, under the sampler’s own ensemble (3), which is what the theorem constrains ($\bar{\kappa} = 3.26$, bound $5.34\sigma^2$; six seeds). Both are proven lower bounds; the second is the tighter, accounting for 81% of the directly measured single-pass floor against 58%. The measured anatomy of the remainder (in-band profile sag) is in the text, as is the tube width at which the comparison expires. Open circles: the interventional test, read under (3) on the same axis as the star — five training protocols move the measured ceiling from 3.2 to 5.8; the measured floor falls from $6.4\text{--}7.2\sigma^2$ (low $\bar{\kappa}$) into the $3.0\text{--}3.4\sigma^2$ range (high $\bar{\kappa}$), with protocol scatter, and every point stays strictly above its bound, the tightest at 78% of its measured floor.

normal width $\sigma_0 \approx 0.072$, thin against its tangent scale ≈ 0.41): with the identical extraction protocol and per-point k NN-PCA normal frames, $\bar{\kappa}$ falls $2.0 \rightarrow 1.6$ as $\sqrt{w}$ runs $0.071 \rightarrow 0.231$ (21%), tightly fit by $\bar{\kappa}(w) \approx 2.2 - 2.7\sqrt{w}$ (max residual 0.017, two seeds). The degradation is therefore not a synthetic-tube artifact. Constants are geometry-specific. Over this range the points are well described by the two-parameter fit $\bar{\kappa}(w) \approx 4.4 - 11.6\sqrt{w}$, an *empirical fit, not a derived form*; nothing below depends on the choice, since refitting linear, logarithmic, or quadratic, or replacing the fit with a fit-free monotone interpolation of the raw points, moves the resulting fixed point by at most 10% (the computation uses the fit only as an interpolant inside the measured range).

Under the law the theorem constrains. The probe above reads the ceiling off the forward process; (3) asks for it under the sampler’s own ensemble. Re-running the identical protocol with both probes on the same networks (`poolwidth_probe_samplerlaw.py`; five widths, three seeds each) reproduces every published forward-process ceiling to three decimals, and gives a sampler-law ceiling 0.830 to 0.841 of them. That band is narrow but not monotone in σ , unlike the $t_0 = 10^{-4}$ probe of Section 5, whose ratio rises from 0.827 to 0.868. The two differ by more than truncation: `poolwidth_probe.py` performs an extra network construction before building the one it trains, which advances the RNG, so at the same nominal seed the two scripts train *different* networks — a 0.3–0.7 shift in $\bar{\kappa}$, larger than any truncation effect. We replicate that construction here so the published ceilings reproduce exactly, average three seeds per width rather than one, and therefore read a constant ≈ 0.83 and

claim no direction; the single-seed trend in §5 should not be read as one. The degradation law is not merely rescaled but *flatter*, $\bar{\kappa}(w) \approx 3.80 - 10.25\sqrt{w}$ against $4.54 - 12.17\sqrt{w}$ (max residual 0.031 and 0.047), so no single factor converts one into the other. Propagated through the same closure, the deterministic fixed point moves from $5.8\sigma^2$ to $10.0\sigma^2$.

Step 2, the self-consistent closure. Feeding the degradation back, $v \leftarrow \Phi_{\text{det}}(w(v), \bar{\kappa}(w(v)))$, moves the fixed point from 4.5 to $5.6\text{--}7.1\sigma^2$ at $t_0 = 10^{-4}$ (the range across the three probe seeds' fits), with the resulting pool width $\sqrt{w^*} \approx 0.091\text{--}0.101$ *inside* the measured support (interpolation, not extrapolation). The closure assumes the degradation depends on the pool only through its width w , while the recursion's actual pools are two-component mixtures (clean tube plus fattened returns); testing this directly by training on the fixed point's own mixture composition gives $\bar{\kappa}^{\text{mix}} = 3.37 \pm 0.15$ against the single-tube prediction 3.34, so the width-only reduction survives its first direct test. This closure shrinks the unexplained factor from 2.4 to 1.5–2.0 — a quarter to a half of the logarithmic gap.

Step 3, the measured remainder. For what remains we first say what it is *not*: iterating the map with the measured generation-to-generation fluctuations gives a Jensen bias of only $0.006\sigma^2$, and applying the measured $t \geq 0.1$ slope-tracking ratio (0.95–0.97) to the whole profile moves the fixed point by only $+0.3\text{--}0.4\sigma^2$. The remainder is then *measured*: the profile-sag and ensemble-displacement channels of the floor anatomy above, re-measured across six pool widths spanning the fixed point, lift the per-pass map to its directly measured values, and iterating that measured map gives a fixed point of $9.9\sigma^2$ against the measured $10.8\text{--}13.2\sigma^2$ — the once factor-2.4 discrepancy is now 1.1–1.3, at the level of protocol differences (mixture pools, per-generation sample counts, seed scatter). What stays underived are the constants themselves: $\bar{\kappa}$, its degradation law, and the saturation profile.

Two steps, not three. Of the two channels Step 3 applies, ensemble displacement is the probe mismatch of Section 5: under (3) it is not a correction to the bound but part of the hypothesis. Measuring $\bar{\kappa}$ under the sampler's own law therefore reaches $10.0\sigma^2$ in Step 1 alone, against the $9.9\sigma^2$ Step 3 reaches by applying it afterwards — the same destination by a shorter route, and to 1% rather than to the level of protocol differences we had expected to have to invoke. The residual factor of 1.1–1.3 against the measured $10.8\text{--}13.2\sigma^2$ is unchanged in size and character; what changes is that one of the two channels no longer has to be carried.

6 The anchoring fix

Corollary 2 says the singular regime cannot be won by scaling the model or tuning the schedule. It can be won by *anchoring the recursion to data*.

The operator. For each point, find its k nearest neighbors in the reference; take the local PCA frame (top- k_{dim} directions tangent, remainder normal); rescale the point's normal coordinates so that the pool's local normal second moments equal the *reference's* local normal second moments in the *same estimated frames*. The match is *bidirectional* (it contracts fat directions and inflates thin ones) and *per-axis*. Both properties are forced by the theory: the sampler's error is signed (fat from injection, thin from posterior-mean sharpening; Proposition 1 in Appendix D) and anisotropic whenever the normal spectrum straddles $s^2(\hat{t})$, so any scalar or one-sided matcher mis-corrects some directions (*we observe*: a scalar matcher leaves MNIST-latent thin ratios at 0.62; the per-axis matcher achieves $|\text{thin} - 1| \leq 0.012$).

Lemma 4 (shared-frame cancellation; *we prove* to first order). *Let both reference and pool normal energies be measured in the same $k\text{NN-PCA}$ frames estimated from the reference. Then the matched*




Figure 4: **Where the anchor sits.** Each generation, the pool (and, for the jump pipeline, the sampler output) is passed through a bidirectional per-axis local moment match against a *stale* reference of $m = 2,000$ real samples drawn once at $g = 0$. No provenance labels are needed; the reference is never refreshed.

pool satisfies, per normal direction, $v(\text{matched}) = \hat{\sigma}_{\text{ref}}^2 + O_p(m^{-1/2}) + \Delta_{\text{leak}}$, where the frame-error leakage Δ_{leak} afflicts pool and reference identically at first order in the frame perturbation and cancels; standard local-PCA rates [15, 16] enter only at second order. Every quantity in the bound is intrinsic (k, τ, σ, m); the ambient dimension d appears nowhere.

How to read the theorem. The collapse of Theorem 1 happens because generation g inherits generation $g-1$'s width: the map F_g has a v_{g-1} inside it, so errors compound. The matcher overwrites the pool's normal width with a fixed number (the reference's) *before* training, so the input to F_g no longer depends on v_{g-1} . The recursion becomes a constant sequence, there is nothing left to compound, and the critical-fraction threshold of Theorem 1 disappears entirely. The only error left is the one-time cost of estimating the reference width from a finite sample, and that does not grow with generations.

Theorem 3 (anchoring makes the recursion memoryless; *we prove* given Lemma 4 and Assumption 1). *Insert the matcher after pooling (and after sampling, for the jump pipeline). Write $\varepsilon_m := O_p(m^{-1/2}) + \Delta_{\text{leak}}$ for the matcher error of Lemma 4, a quantity that does not grow with g . Then the training input has normal variance $\hat{\sigma}_{\text{ref}}^2 + \varepsilon_m$ regardless of v_{g-1} : the map (2) loses its v -argument,*

$$v_g = F_g(\hat{\sigma}_{\text{ref}}^2 + \varepsilon_m),$$

*a constant sequence up to the non-compounding ε_m . Consequently: no fixed-point equation, **no critical fraction, and stability for every $\lambda > 0$** ; and with S_{trunc} the dynamics reduce to truncation-only, restoring the conclusion of Khelifa et al. [1]'s Theorem 4.1 in the singular regime (the data remains singular; what anchoring removes is the compounding term that their hypotheses excluded).*

Why this is not a replay buffer. Replay, data-accumulation, and data-refresh remedies [4, 8] fight collapse by *re-inserting real examples* into the pool: they preserve *samples*. The anchor preserves something weaker and more portable, the *local normal moments* (per-axis widths) of a small stale reference, and imposes them on *whatever the model currently generates*. It never adds a




Figure 5: **Head-to-head vs. the state-of-the-art remedy (*we observe*).** Ring manifold, $\sigma = 0.05$, $\lambda = \frac{1}{2}$, 20 generations, 5 seeds (mean $\pm$ s.d.). Red: annealed truncation $t_0: 0.05 \rightarrow 10^{-4}$, no anchor [1]; it rises to $47\sigma^2$, then thins only to $9.8\sigma^2 \pm 1.8$. Blue/green: the same recursion with the anchor (standard and jump samplers): $1.01\sigma^2 \pm 0.02$, flat for 20 generations.

reference point to the training pool, uses no provenance labels, and reuses the same $m=2,000$ points for every generation. The distinction is not cosmetic. A replay buffer of size m still lets the $(1 - \lambda)$ synthetic majority set the pool’s fine-scale statistics, so the floor of Section 5 still governs that majority; anchoring instead *overwrites* the fine-scale statistics of the majority itself, which is why it removes the floor rather than diluting it. In the language of the law (Theorem 1): replay changes λ ; anchoring changes the response map F_g . We test this at the anchor’s exact information budget (*we observe*): inserting the anchor’s own stale reference ($m = 2,000$ real points, drawn once at $g = 0$, never refreshed) *into the pool* every generation of the head-to-head protocol — raising the effective real fraction to $\lambda_{\text{eff}} = 2/3$, with synthetic data now a $1/3$ *minority* — moves the 20-generation plateau only from 9.8 ± 1.8 to $8.5 \pm 1.1 \sigma^2$ (5 seeds, same annealed schedule), the small dilution the law predicts for a larger λ ; the anchor, using the same 2,000 points as a moment reference instead, holds $1.01 \pm 0.02 \sigma^2$. Replay dilutes the floor; anchoring removes it.

Attribution, stated plainly. In our pipeline the *matcher* is the floor-breaker; the Tweedie jump is a *raw-error reducer*. The jump’s raw output is itself fat (7 to $20 \times \sigma^2$ across settings); its value is a $\approx 2\times$ smaller one-shot raw error and $\approx 3\times$ smaller unanchored plateau (21 vs $68 \sigma^2$ in raw normal variance, at $\lambda = \frac{1}{2}$ and matched moderate truncation $t_0 = \hat{t} = 0.05$; the standard plateau includes the truncation remainder $s^2(0.05)/\lambda \approx 39 \sigma^2$ that the jump removes, and the jump plateau is itself predicted by Theorem 1 as $e_{\text{jump}}^2/\lambda$), plus smaller post-match tails. Post-anchor, the standard and jump samplers converge to the same variance; we report both.

Dimension scaling (*we observe*). Lemma 4’s d -independence is a postdiction: across ambient $d = 8 \rightarrow 128$ (curved 2-manifolds, fixed tube), the *raw* one-shot floors grow steeply ($\sim d^{0.76}$ standard, $\sim d^{0.95}$ jump, the collapse itself), while the *matched* floors are flat ($\sim d^{0.1}$), and anchored recursions are flat for 8 generations at both $d = 8$ and $d = 128$.




Figure 6: **Real pixels (*we observe*)**. Self-consuming recursion on 8×8 MNIST (64-dimensional pixel space), $\lambda = \frac{1}{2}$, 8 generations, 3 seeds. Off-manifold distance = mean nearest-neighbor distance to a held-out real test set. Unanchored drifts to 3.06 ($2.3\times$ the true-data baseline 1.34); anchored holds 1.66 ($1.24\times$), closing $\sim 81\%$ of the gap. The mild upward creep of the anchored arm reflects the local-chart approximation on a genuinely curved manifold.

7 Experiments

All core experiments (ring, curved manifolds, MNIST, law identification) are CPU-only (≤ 12 threads) and use small MLPs (3 hidden layers, width 256, SiLU, sinusoidal time embedding) with denoising score matching and geometric reverse-time grids; the CIFAR-10 scale-up alone uses a ~ 7 M-parameter convolutional denoiser on a single consumer GPU (mixed precision). Every long run streams to an append-only log and is exactly resumable. Full protocols, seeds, and a script-to-claim map are in the released code (Section 9).

Protocols. *Ring*: radius 2.5, $\sigma = 0.05$ ($\sigma^2 = 0.0025$), 4,000 samples/generation, 2,000 training steps. *Head-to-head* (Figure 5): $G = 20$, $\lambda = \frac{1}{2}$, 5 seeds; arm A anneals $t_0: 0.05 \rightarrow 10^{-4}$ geometrically (the remedy of 1); arms B/C add the anchor to the standard/jump samplers. *Two-horned ablation* (Figure 2): fixed $t_0 = 10^{-4}$, noise on/zeroed, $G = 15$, 3 seeds. *Replay control*: head-to-head protocol with the anchor’s $m=2,000$ stale reference inserted into the pool each generation instead of used as a moment reference ($G = 20$, 5 seeds). *Critical fraction*: fixed $t_0 = 10^{-4}$, $\lambda \in \{0.02, 0.25\}$, 3 seeds each. *Pixel MNIST* (Figure 6): 8×8 pixels, $G = 8$, 3 seeds, local matcher with $k = 96$, $k_{\text{dim}} = 12$. *Law identification*: fixed- t_0 recursions at $t_0 \in \{0.02, 0.05\}$, $\lambda \in \{0.25, 0.5, 0.75\}$; the affine law fitted at one t_0 predicts fixed points at the other within 1 to 13% (Assumption 1 validation); the direct probe slope of e^2 is small and saturating ($+0.13 \rightarrow +0.05$ over the upper measured range), the net of a measured deficit channel -0.68 and injection channel $+0.73$, and the recursion’s own convergence rates independently give a slope falling from $+0.6$ (thin fixed-point pools, $\lambda = 0.75$) to ≈ 0 (fat, $\lambda = 0.25$) (Assumption 2 validation); a refinement using the probe’s own capacity measurements splits the injection channel further (Appendix F).




Figure 7: **CIFAR-10 pixel recursion (*we observe*)**. (a) Off-manifold distance (mean nearest-neighbor distance to held-out real CIFAR-10) across 8 self-consuming generations, 3 seeds (mean $\pm$ s.d.). The two arms stabilize on opposite sides of the real-vs-real baseline (20.15): unanchored above (+2.3, the predicted fattening), anchored below (−3.9, mild contraction toward the manifold); the sign of the split holds in every seed. Because both arms deviate from the baseline, panel (a) alone establishes stability, not superiority; the ranking evidence is (b): three nearest-neighbor metrics relative to a real-vs-real baseline (dotted). The anchor keeps *coverage* of the real data at 1.00 $\times$ (no real modes dropped, ruling out mode collapse) while the unanchored arm loses coverage (1.16 $\times$); the anchor’s 0.75 $\times$ diversity is a bounded, non-compounding contraction cost (identical at generations 4 and 7).

The critical fraction, located and exhibited (*we observe*). The probe’s saturated slope, $c_\infty \approx 0.03\text{--}0.05$, puts the critical fraction at $\lambda^* = c_\infty / (1 + c_\infty) \approx 0.03\text{--}0.05$ for this pipeline — small, so the collapse is self-limiting at any practically relevant fresh-data fraction, and the theorem’s content is that $\lambda^* > 0$ at all. A recursion at $\lambda = 0.02$ (at or below the estimate; fixed $t_0 = 10^{-4}$, 3 seeds, $G = 25$) exhibits the subcritical regime directly: the width climbs monotonically through $69 \pm 16 \sigma^2$ at $g = 24$, roughly 10 $\times$ its early value, with no sign of flattening — while a matched $\lambda = 0.25$ control plateaus at $16.5 \pm 2.2 \sigma^2$ by generation ~ 5 . (A finite horizon cannot distinguish true divergence from a plateau far above it, so we claim the contrast, not divergence itself.) The supercritical plateaus track the law’s e^2/λ scaling: the implied per-pass excess $e^2(w^*) = \lambda(v_\infty - \sigma^2)$ is 3.9–5.0 σ^2 across $\lambda \in \{0.25, 0.5, 2/3\}$ (the last from the replay control of Section 6), sitting between the envelope floor and the directly measured per-pass values, as the width-dependence of the measured map requires.

Scaling to CIFAR-10 pixels (*we observe*). To test whether the mechanism survives at realistic dimension, we ran the same self-consuming recursion in raw CIFAR-10 pixel space ($32 \times 32 \times 3 = 3,072$ dimensions) with a small convolutional denoiser ($\sim 7\text{M}$ parameters, mixed precision, single consumer GPU), $\lambda = \frac{1}{2}$, $G = 8$, three seeds (Figure 7a). Both arms are stable, and they stabilize on opposite sides of the real-vs-real baseline (20.15): the unanchored recursion settles *above* it (off-manifold distance 22.5 ± 0.5 , a deviation of +2.3, the fattening the theory predicts, same sign in every seed), while the anchored recursion sits *below* it (16.2 ± 0.3 , a deviation of −3.9, mild contraction toward the manifold reflecting the anchor’s known diversity cost). We are explicit about what panel (a)

Table 1: Headline numbers (mean $\pm$ s.d. over seeds; variance in units of σ^2 except MNIST and CIFAR-10, which report off-manifold distance).

Experiment	Arm	Final value	Seeds
Head-to-head, $G=20$	annealed truncation [1]	9.84 ± 1.85	5
	standard sampler + anchor	1.01 ± 0.02	5
	jump sampler + anchor	1.01 ± 0.02	5
	annealed + replay ($m=2,000$ in-pool)	8.5 ± 1.1	5
Critical fraction, $t_0=10^{-4}$	$\lambda = 0.02$ (subcritical)	69 ± 15 at g_{24} , rising	3
	$\lambda = 0.25$ (control)	16.5 ± 2.2 (plateau)	3
Ablation, fixed $t_0=10^{-4}$	stochastic (noise on)	10.8 to 13.2	3
	deterministic (noise zeroed)	0 (point collapse)	3
Pixel MNIST, $G=8$	unanchored	3.06	3
	anchored	1.66	3
	true-data baseline	1.34	n/a
Pixel CIFAR-10, $G=8$	unanchored	22.5 ± 0.5	3
	anchored	16.2 ± 0.3	3
	true-data baseline	20.15	n/a
	anchored coverage / diversity	$1.00\times$ / $0.75\times$	3
Floor law	envelope at $\bar{\kappa}$ under p_t / under (3)	3.8 / $5.5\sigma^2$ (58% / 81% of measured)	6
	interventional test, $\bar{\kappa}: 3.2 \rightarrow 5.8$	floor $\{6.4\text{--}7.2\} \rightarrow \{3.0\text{--}3.4\}$, 0/5 violated	3

can and cannot show: both arms deviate from the baseline, so the distance metric alone ranks neither arm; it establishes that the anchored dynamics are stable and non-compounding, not that its samples are closer to the data distribution. The ranking evidence is Figure 7b: against a real-vs-real baseline, the anchored arm keeps *coverage* of the real data at $1.00\times$ (no real modes dropped, ruling out mode collapse) while the unanchored arm loses coverage ($1.16\times$); the anchored arm’s reduced intra-set diversity ($0.75\times$) is a bounded cost, identical at generations 4 and 7 (non-compounding), whereas the unanchored arm’s coverage loss is the collapse itself. At this model budget the samples are blurry; this is a check on the collapse *dynamics* and a ruling-out of mode collapse, not a demonstration of perceptual quality. The unanchored arm reaches a stable displaced fixed point rather than runaway collapse, exactly as Theorem 1 predicts at $\lambda = \frac{1}{2}$, where fresh-data injection pins a finite plateau.

8 Limitations and scope

Everything in the *we model* tier is collected here. None of it touches the paper’s unconditional core: the closed-form floor (Theorem 2) and the $\Omega(\sqrt{J})$ capacity requirement (Corollary 2) hold given only the slope cap $\bar{\kappa} < \infty$, and the memorylessness theorem (Theorem 3) holds given the frame-cancellation lemma; the items below concern how tightly our *measured* constants pin the recursion’s plateau, not whether the floor exists. Item (8) is of a different kind: it delimits which architectures satisfy the slope-cap hypothesis in a binding way, without bearing on the theorem itself. (1) Assumptions 1 and 2 are measured, cross-validated laws about the trained pipeline, not derived facts; we present them as such, in the spirit of measured critical exponents. (2) The recursion floor constant is accounted for empirically, not derived: our closed forms are exact for a single pass and lower bounds for the recursion; feeding the measured capacity degradation $\bar{\kappa}(w)$

back self-consistently closes a quarter to a half (in log) of the once factor-2.4 gap; measuring that degradation under the law (3) constrains, rather than under p_t , closes it to 1.1–1.3 $\times$ with no further correction, and the per-pass map’s one remaining measured channel (in-band profile sag) accounts for the rest (Section 5), at the level of protocol differences. Separately, the *numerical* envelope comparison accounts for 77–85% of the measured floor at $\sigma = 0.05$ and 90–97% at $\sigma = 0.10$, and exceeds it in three of six seeds at $\sigma = 0.13$ by at most 3.4%. That is inside the relative error term $O(\sigma^2/\rho^2) = 4.2\%$ that Theorem 2 carries at that width, and it also disappears under the curved-geometry bound these ring networks actually obey (Corollary 1), which no seed at any width exceeds; no validity range is restricted. Every step of this accounting is a measurement; none of the constants ($\bar{\kappa}$, its degradation law — an empirical fit replicated across four synthetic geometries and real MNIST 8×8 pixel data — or the saturation profile) is derived. (3) The residual-field ($\sim 30\%$ single-pass) channel is bounded in a block-correlation model (Appendix E); negative association across blocks is the one direction that could weaken that particular bound (the closed-form floor of Theorem 2 does not depend on it). (4) Curvature enters in two distinct places and we now treat them differently. In the *variance flow* it is exact: (5) carries it as $\text{Cov}(h, \Delta h)$ with no error term, and Corollary 1 bounds its effect on the floor at 2–3% for our geometries, with the caveat on rare small-radius stragglers recorded there. In the *flat-tube idealization* of the fiber marginal — Gaussian and decoupled from tangential position — curvature is still carried at $O(s^2/\tau)$ and we do not compute those constants; the sphere experiments bound the resulting bias empirically (non-compounding). The two should not be conflated: the first is now a closed form, the second is not. (5) The anchor controls normal *second moments*; tails and tangential statistics improve but are not anchored; on strongly curved data the anchored recursion shows a mild residual creep (Figure 6). (6) A parameter-free one-step prediction of MNIST-latent thin ratios fails quantitatively (it predicts the direction and the regime boundary, not the number); we flag it rather than fit it. (7) Our evidence spans a ring, curved 2-manifolds swept across ambient dimension $d = 8 \rightarrow 128$, six (t_0, λ) law-identification cells, the two-horned ablation, a replay-buffer control and a subcritical-fraction run, pixel-space MNIST, and a preliminary CIFAR-10 pixel-space recursion (3,072 dimensions, three seeds), all chosen so that every quantity the theory names is directly measurable. The CIFAR-10 run confirms the collapse *dynamics* and rules out mode collapse at scale (Figure 7), but at this model budget the samples are blurry and we do not claim a perceptual-quality result; we do not test latent-diffusion pipelines or larger models. The theory predicts the mechanism transfers, because the floor depends only on the *local* normal slope and not on the ambient dimension (Lemma 4, and the measured d -independence of the matched floors); confirming this in a large latent-diffusion system is the clearest next step and is left as future work. (8) **Architectures that factor out the singular normal factor.** The floor’s hypothesis is that the learned normal slope is bounded, $\hat{\kappa} \leq \bar{\kappa}$, which binds when the network must itself represent the steep normal response. It does *not* bind for parameterizations that build the singular factor into the sampler by construction, for instance a score of the form $-(x - \hat{x})/h(t)$ with a learned projection-like map $\hat{x}$, as in the SiLD construction of Huang et al. [12]: there the network supplies $\hat{x}$ rather than the slope, and $\hat{\kappa}$ is not limited by network capacity. Theorem 2 is silent about such models. The experiments here all use ordinary ε -prediction, for which the hypothesis does bind. We thank Wei Huang for this observation. The finite-capacity question does not disappear in that setting, it relocates: writing $\hat{x} = \Pi_M(x) + \delta(x)$, the score error carries a term of order $\delta(x)/h(t)$, so projection error is amplified by the same singular factor as $t \rightarrow 0$. Whether that amplification admits a floor of the same $\bar{\kappa}^{-2}$ character, with $\mathbb{E}[\delta^2]$ in place of the slope deficit, is open and is the clearest theoretical extension of this work.

9 Reproducibility

Code, result logs, and figures are public at

<https://github.com/Naman84789/model-collapse-manifolds>

Every claim maps to a script, an append-only result log, and a figure in the released code (ring, ablation, replay and critical-fraction controls, MNIST, floor law, law identification; ~ 3 CPU-hours total). The result figures regenerate from logged data in seconds. A README lists the run order, the load-bearing scripts, and the headline numbers a reproduction must match. The released code also includes an adversarial verification script for Theorem 2: it re-derives the band crossing symbolically (independent of the manuscript’s derivation), integrates the limit ODE with an independent solver, confirms the limiting cases $\rho \rightarrow 0$, $\rho \rightarrow 1^-$, $\bar{\kappa} \rightarrow \infty$, and $\lambda \rightarrow 1$, verifies that the boundary layer contracts at the predicted quadratic rate, and attacks the comparison bound with 200 randomized slope profiles including overshooting ones; all checks pass.

Acknowledgments

Numerical experiments, figure preparation, and portions of the derivations were carried out with the assistance of an AI system (Claude, Anthropic), used as a research tool under the author’s direction; the author reviewed all mathematics and takes full responsibility for every claim.

References

- [1] Naïl B. Khelifa, Richard E. Turner, and Ramji Venkataramanan. Recursively trained diffusion models: Limiting collapse distribution and spectral characterization. *arXiv preprint arXiv:2606.13796*, 2026.
- [2] Ilia Shumailov, Zakhar Shumaylov, Yiren Zhao, Nicolas Papernot, Ross Anderson, and Yarin Gal. AI models collapse when trained on recursively generated data. *Nature*, 631:755–759, 2024.
- [3] Sina Alemohammad, Josue Casco-Rodriguez, Lorenzo Luzi, Ahmed Imtiaz Humayun, Hossein Babaei, Daniel LeJeune, Ali Siahkoobi, and Richard G. Baraniuk. Self-consuming generative models go MAD. In *International Conference on Learning Representations (ICLR)*, 2024.
- [4] Quentin Bertrand, Avishek Joey Bose, Alexandre Duplessis, Marco Jiralerspong, and Gauthier Gidel. On the stability of iterative retraining of generative models on their own data. In *International Conference on Learning Representations (ICLR)*, 2024.
- [5] Elvis Dohmatob, Yunzhen Feng, Pu Yang, François Charton, and Julia Kempe. A tale of tails: Model collapse as a change of scaling laws. In *International Conference on Machine Learning (ICML)*, 2024.
- [6] Charles Fefferman, Sanjoy Mitter, and Hariharan Narayanan. Testing the manifold hypothesis. *Journal of the American Mathematical Society*, 29(4):983–1049, 2016.
- [7] Jakiw Pidstrigach. Score-based generative models detect manifolds. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [8] Matthias Gerstgrasser, Rylan Schaeffer, Apratim Dey, Rafael Rafailov, Henry Sleight, John Hughes, Tomasz Korbak, Rajashree Agrawal, Dhruv Pai, Andrey Gromov, Daniel A. Roberts,

- Diya Yang, David L. Donoho, and Sanmi Koyejo. Is model collapse inevitable? breaking the curse of recursion by accumulating real and synthetic data. *arXiv preprint arXiv:2404.01413*, 2024.
- [9] Nail B. Khelifa, Richard E. Turner, and Ramji Venkataramanan. Quantifying error propagation and model collapse in diffusion models. *arXiv preprint arXiv:2602.16601*, 2026.
 - [10] Valentin De Bortoli. Convergence of denoising diffusion models under the manifold hypothesis. *Transactions on Machine Learning Research (TMLR)*, 2022.
 - [11] Eliot Beyler and Francis Bach. Optimal denoising in score-based generative models: The role of data regularity. *arXiv preprint arXiv:2503.12966*, 2025.
 - [12] Wei Huang, Andi Han, Mingyuan Bai, Huanjian Zhou, Qixin Zhang, Taiji Suzuki, and Kenji Fukumizu. Provably learning diffusion models under the manifold hypothesis: Collapse and refine. *arXiv preprint arXiv:2605.20235*, 2026.
 - [13] Zichen Liu, Wei Zhang, and Tiejun Li. Improving the euclidean diffusion generation of manifold data by mitigating score function singularity. *arXiv preprint arXiv:2505.09922*, 2025.
 - [14] Bradley Efron. Tweedie’s formula and selection bias. *Journal of the American Statistical Association*, 106(496):1602–1614, 2011.
 - [15] Eddie Aamari and Clément Levrard. Nonasymptotic rates for manifold, tangent space and curvature estimation. *The Annals of Statistics*, 47(1):177–204, 2019.
 - [16] Amit Singer and Hau-Tieng Wu. Vector diffusion maps and the connection Laplacian. *Communications on Pure and Applied Mathematics*, 65(8):1067–1144, 2012.
 - [17] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
 - [18] Yang Song, Jascha Sohl-Dickstein, Diederik P. Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. In *International Conference on Learning Representations (ICLR)*, 2021.
 - [19] Iskander Azangulov, George Deligiannidis, and Judith Rousseau. Convergence of diffusion models under the manifold hypothesis in high-dimensions. *arXiv preprint arXiv:2409.18804*, 2025.

A Proof of Theorem 1

Write $h(v) := T(v) - v$ with T as in (2) (we reserve g for the generation index); h is concave (composition of an affine map with the concave nondecreasing $w \mapsto w + e^2(w)$, minus a linear term), with $h(0) = T(0) = \lambda\sigma^2 + e^2(\lambda\sigma^2) > 0$ and one-sided derivative $h'(v) = (1 - \lambda)(1 + e^{2'}(w(v))) - 1$, nonincreasing in v with limit $(1 - \lambda)(1 + c_\infty) - 1 < 0$ when $\lambda > \lambda^*$.

(i) Existence, uniqueness. A concave function positive at 0 whose derivative is eventually $\leq -\delta < 0$ tends to $-\infty$, so h has a zero; concavity plus $h(0) > 0$ forbids two downcrossings, so the zero v_∞ is unique, with $T(v) > v$ for $v < v_\infty$ and $T(v) < v$ for $v > v_\infty$.

Contraction above. For $v \geq v_\infty$, monotonicity and concavity give $0 \leq T(v) - v_\infty \leq \rho_{\text{loc}}(v - v_\infty)$ with $\rho_{\text{loc}} = T'(v_\infty^+) < 1$. The strict inequality needs no extra hypothesis: if $\rho_{\text{loc}} = 1$ then $h'(v_\infty) = 0$, making the concave h maximal at v_∞ , whence $h(v_\infty) \geq h(0) = T(0) > 0$, contradicting $h(v_\infty) = 0$. So the crossing is always strict and the rate genuinely geometric.

Convergence under vanishing perturbations. Let $v_g = T(v_{g-1}) + s_g^2$, $s_g^2 \rightarrow 0$. While iterates stay $\geq v_\infty$, $x_g := v_g - v_\infty \leq \rho_{\text{loc}}x_{g-1} + s_g^2 \rightarrow 0$ by the standard perturbed contraction argument (split $\sum_j \rho_{\text{loc}}^{g-j} s_j^2$ at $j = g/2$). If an iterate drops below v_∞ : define the unperturbed comparison $u_g = T(u_{g-1})$ started there; then $u_g \uparrow v_\infty$ (monotone, bounded, $T(v) > v$ below v_∞) and $v_g \geq u_g$ by induction ($s_g^2 \geq 0$, T monotone), so $\liminf v_g \geq v_\infty$, while $v_g < v_\infty + s_g^2$ gives $\limsup v_g \leq v_\infty$. Alternating regimes inherit both barriers. Linearizing T at v_∞ gives the local rate.

(ii) Divergence. If $(1 - \lambda)(1 + c_\infty) > 1$ then, since $e^{2'} \geq c_\infty$ everywhere (the concave derivative decreases to its limit), $h' \geq (1 - \lambda)(1 + c_\infty) - 1 > 0$, so h is nondecreasing and $h(v) \geq h(0) > 0$ for all v : after G generations $v_G \geq v_0 + GT(0) \rightarrow \infty$, regardless of the schedule ($s_g^2 \geq 0$ only helps).

(iii) Substituting $e^2(w) = e_0^2 + cw$ makes T affine with slope $(1 + c)(1 - \lambda)$; solve the fixed point. $\square$

B The variance ODE

Lemma 1. One Euler step of the reverse scheme is $x' = x + (\frac{1}{2}x - \hat{\varepsilon}(x, t)/s)dt + \sqrt{dt}\xi$ with ξ independent of x . Write $f(x) = \frac{1}{2}x - \hat{\varepsilon}(x, t)/s$. By bilinearity of the covariance and independence of the injected noise,

$$\text{Var}(x') = \text{Var}(x) + 2dt \text{Cov}(x, f(x)) + dt^2 \text{Var}(f(x)) + dt = V + \left[V - \frac{2}{s} \text{Cov}(x, \hat{\varepsilon}(x)) + 1 \right] dt + O(dt^2).$$

No property of $\hat{\varepsilon}$ beyond square-integrability is used, so this holds for a nonlinear estimator; the $+dt$ is the injected noise and the $O(dt^2)$ term is the finite variance of the drift. Divide by dt , let $dt \rightarrow 0$, and substitute $\text{Cov}(x, \hat{\varepsilon}(x)) = \hat{\kappa} \text{Var}(x)$ from (3) to get (4).

Independent route, no discretisation. For the continuous reverse SDE $dX = f(X)d\tau + dW$, Itô gives $d(X^2) = 2Xf(X)d\tau + 2XdW + d\tau$, so $\mathbb{E}[X^2]' = 2\mathbb{E}[Xf(X)] + 1$ while $\mathbb{E}[X]' = \mathbb{E}[f(X)]$; subtracting, $\text{Var}' = 2\text{Cov}(X, f(X)) + 1$, which is (4). With $\hat{\kappa} = \kappa = s/\gamma^2$: $\kappa/s = 1/\gamma^2$ and $V = \gamma^2$ solves the equation exactly (direct check), confirming the bookkeeping. An intercept in $\hat{\varepsilon}$ moves only the mean. $\square$

The curved case, (5). Let h be the signed distance to the manifold, so $|\nabla h| = 1$ and $n := \nabla h$. For $dX = f(X)d\tau + dW$ in $\mathbb{R}^D$, Itô gives $dh = (\langle n, f \rangle + \frac{1}{2}\Delta h)d\tau + \langle n, dW \rangle$, and $\langle n, dW \rangle$ has quadratic variation $d\tau$ because $|n| = 1$. The same subtraction as above, with drift $B = \langle n, f \rangle + \frac{1}{2}\Delta h$, gives $\text{Var}'(h) = 2\text{Cov}(h, B) + 1$ exactly — no expansion, and no property of $\hat{\varepsilon}$ beyond square-integrability. With $f = \frac{1}{2}x - \hat{\varepsilon}/s$ we have $\langle n, f \rangle = \frac{1}{2}\langle n, x \rangle - \varepsilon_n/s$ where $\varepsilon_n := \langle n, \hat{\varepsilon} \rangle$, so $2\text{Cov}(h, \langle n, f \rangle) = \text{Cov}(h, \langle n, x \rangle) - 2\hat{\kappa}V/s$; adding $2\text{Cov}(h, \frac{1}{2}\Delta h)$ and writing $\text{Cov}(h, \langle n, x \rangle) = V + [\text{Cov}(h, \langle n, x \rangle) - V]$ gives (5). Two checks on the bracket: for a plane through the origin, $h = x_D$ and $\langle n, x \rangle = x_D = h$;

for a sphere about the origin, $h = r - \alpha R$ and $\langle n, x \rangle = r = h + \alpha R$. In both cases $\text{Cov}(h, \langle n, x \rangle) = V$ and the bracket vanishes. For the plane $\Delta h = 0$ as well, recovering (4).

The mean equation, (7). Taking expectations in the same Itô expansion, the martingale part drops and $\mathbb{E}[h]' = \mathbb{E}[\langle n, f \rangle] + \frac{1}{2}\mathbb{E}[\Delta h]$ in reverse time. With $f = \frac{1}{2}x - \hat{\varepsilon}/s$ and $\langle n, x \rangle = h + c$ for a constant c on the two geometries above, $\mathbb{E}[\langle n, f \rangle] = \frac{1}{2}\mu + \frac{1}{2}c - \mathbb{E}[\varepsilon_n]/s$; absorbing the constant into the drift of the contracted centre gives (7). Note the contrast with (6): the variance correction is $\text{Cov}(r, 1/r)$, which vanishes as the tube contracts because r concentrates, whereas the mean correction is $\frac{D-1}{2}\mathbb{E}[1/r] \rightarrow \frac{D-1}{2\alpha R}$, which does not. That is why the mean equation remains a sharp test at every t while the variance equation stops being one.

For (6): $\text{Cov}(r, 1/r) = \mathbb{E}[r \cdot \frac{1}{r}] - \mathbb{E}[r]\mathbb{E}[1/r] = 1 - \mathbb{E}[r]\mathbb{E}[1/r]$, and the two-point form of Remark 2 with $g(r) = 1/r$ gives $\text{Cov}(r, 1/r) = \frac{1}{2}\mathbb{E}[(r - r')(\frac{1}{r} - \frac{1}{r'})] = -\frac{1}{2}\mathbb{E}[(r - r')^2/(rr')] \leq 0$. The same two-point form with $|g'| = (D-1)/r^2 \leq (D-1)/a^2$ on $r \geq a$ gives Corollary 1. $\square$

Remark 2. For an i.i.d. copy X' of X , $\frac{1}{2}\mathbb{E}[(X - X')(g(X) - g(X'))]$ expands to $\mathbb{E}[Xg(X)] - \mathbb{E}[X]\mathbb{E}[g(X)] = \text{Cov}(X, g(X))$. If g is $\bar{\kappa}$ -Lipschitz then $(X - X')(g(X) - g(X')) \leq \bar{\kappa}(X - X')^2$ pointwise, and $\mathbb{E}[(X - X')^2] = 2\text{Var}(X)$, so $\text{Cov}(X, g(X)) \leq \bar{\kappa}\text{Var}(X)$ under any law. $\square$

C Floor lemmas and closed forms

Lemma 2. From (1), $\kappa(t) = s/(a^2\sigma^2 + s^2)$ with $s^2 = 1 - e^{-t}$: $\kappa \rightarrow 0$ as $t \rightarrow 0$ ($s \rightarrow 0$, $\gamma^2 \rightarrow \sigma^2$), $\kappa \approx 1/s$ for $s^2 \gg \sigma^2$, and maximizing over s gives $\kappa_{\max} = 1/(2\sigma)$ at $s = \sigma$ (i.e. $t \approx \sigma^2$), to leading order in σ^2 . Below the band the remaining e-folds are $\int_{t_0}^t 2\kappa/s dt = \int 2 dt/\gamma^2 = 2 \log \frac{\sigma^2 + t_-}{\sigma^2 + t_0} \leq 2 \log(1 + \rho^2/4) \leq \rho^2/2$, using $t_- \approx \bar{\kappa}^2\sigma^4$. $\square$

Lemma 3. Substitute $t = \sigma^2 u$, $V = \sigma^2 W$: $s^2 = \sigma^2 u(1 + O(\sigma^2 u))$, $\gamma^2 = \sigma^2(1 + u)(1 + O(\sigma^2 u))$, so $\kappa = (1/\sigma)\frac{\sqrt{u}}{1+u}(1 + O(\sigma^2 u))$ and $2\hat{\kappa}/s = (2m(u)/\sqrt{u})/\sigma^2$; the ODE becomes $-dW/du = 1 - (2m/\sqrt{u})W + \sigma^2 W(1 + o(1))$. On $u \in [0, U]$ the perturbing terms are $O(\sigma^2 U)$ per unit of W ; the band occupies $u \leq u_+ = O(1/\rho^2)$, giving the stated relative error $O(\sigma^2/\rho^2)$ (the additive gap to g is a factor $W \sim \rho^{-2}$ larger, hence $O(\sigma^2/\rho^4)$). For the boundary: the uncapped equation $-dW/du = 1 - 2W/(1+u)$ has general solution $W = (1+u) + A(1+u)^2$ (direct check), so deviations from the slave solution $W = 1+u$ contract as $((1+u)/(1+u'))^2$ under downward integration from u' to u ; the influence of the $t \gtrsim 1$ boundary layer at band entry is suppressed by this quadratic factor and is absorbed in the same error term. *Numerical check:* the slave solution is invariant to 3×10^{-15} under the integrator; finite- σ floors converge to the limit values from above (e.g. $\rho = 0.39$: 3.953/3.812/3.775 at $\sigma = 0.10/0.05/0.02$ vs. limit 3.769). $\square$

Theorem 2(U). By the scalar comparison principle (equal forcing and data; smaller $\hat{\kappa}$ means larger V), any profile with $\hat{\kappa} \leq \bar{\kappa}$ has $V \geq V_{\hat{\kappa}=\bar{\kappa}}$, one-sided; no lower bound on $\hat{\kappa}$ is used. For the constant envelope, (9) with $m \equiv \rho/2$ has integrating factor $e^{-2\rho\sqrt{u}}$: $\frac{d}{du}[We^{-2\rho\sqrt{u}}] = -e^{-2\rho\sqrt{u}}$, so for data of sub-quadratic growth at $u = \infty$, $W(0^+) = \int_0^\infty e^{-2\rho\sqrt{u}} du = 2 \int_0^\infty xe^{-2\rho x} dx = \frac{1}{2\rho^2}$, exactly, for every $\rho > 0$. $1/(2\rho^2) > 1 \iff \rho < 1/\sqrt{2}$.

Theorem 2(N). The band edges solve $\sqrt{u}/(1+u) = \rho/2$, a quadratic in $x = \sqrt{u}$: $\rho x^2 - 2x + \rho = 0$, $x_\mp = (1 \mp \rho)/\rho$, $q = \sqrt{1 - \rho^2}$ (nonempty iff $\rho < 1$). Above the band the solution is the exact slave $W = 1 + u$; it enters the band with $W(u_+) = 1 + u_+$. Across the band ($m = \rho/2$), the same integrating factor with antiderivative $F(u) = -(2\rho\sqrt{u} + 1)e^{-2\rho\sqrt{u}}/(2\rho^2)$ gives

$$W(u_-) = (1 + u_+)e^{-2\rho(x_+ - x_-)} + \frac{(2\rho x_- + 1) - (2\rho x_+ + 1)e^{-2\rho(x_+ - x_-)}}{2\rho^2},$$

and $2\rho x_\mp = 2(1 \mp \rho)$, $2\rho(x_+ - x_-) = 4q$ yield the displayed W_{exit} . Below the band the equation is uncapped with exact general solution $(1 + u) + A(1 + u)^2$; matching at u_- gives $A = (W(u_-) - (1 + u_-))/(1 + u_-)^2$.

$u_-)) / (1 + u_-)^2$ and $g(\rho) = W(0^+) = 1 + A$, which is (11). Limits: as $\rho \rightarrow 1^-$ ($q \rightarrow 0$), $u_{\pm} \rightarrow 1$ and $g \rightarrow 1$; as $\rho \rightarrow 0$ ($q \rightarrow 1$), $u_+ \approx 4/\rho^2$, $u_- \approx \rho^2/4$, and expanding, $g = [8e^{-4} + 1 - 5e^{-4}] / (2\rho^2) + O(1) = (1 + 3e^{-4}) / (2\rho^2) + O(1)$, giving $C_0 = (1 + 3e^{-4})/2$. Positivity $g > 1$ on $(0, 1)$ follows from $W(u_-) > 1 + u_-$ (the band under-contracts relative to the slave). *Numerical check:* (11) and the exact $1/(2\rho^2)$ match direct integration of (9) to $\leq 2.4 \times 10^{-5}$ over $\rho \in [0.1, 1]$ (pure Euler residual).

Both bounds apply to the fully nonlinear estimator with no extension step, since (4) is exact and $\hat{\kappa}$ is defined under the sampler's own law (Lemma 1). Lemma 3 transfers the limit values to finite σ with the stated error. $\square$

Corollary 2 is Theorem 2(U) read at fixed $\bar{\kappa}$ with $\sigma \rightarrow 0$; the capacity bound inverts $1/(2\rho^2) \leq C$. $\square$

D Posterior-mean sharpening and per-axis necessity

Proposition 1 (Tweedie sharpening; *we prove*). *In the Gaussian normal coordinate with axis variance σ_{ax}^2 , the jump with the exact predictor outputs the posterior mean, with*

$$\text{Var}(\hat{x}_0) = \frac{a^2 \sigma_{\text{ax}}^4}{\gamma_{\text{ax}}^2(\hat{t})}, \quad \frac{\text{Var}(\hat{x}_0)}{\sigma_{\text{ax}}^2} = \frac{a^2 \sigma_{\text{ax}}^2}{a^2 \sigma_{\text{ax}}^2 + s^2(\hat{t})} < 1,$$

and with a trained predictor $\hat{\varepsilon} = \varepsilon^* - \tilde{\delta}$, $v_{\text{jump,ax}} = a^2 \sigma_{\text{ax}}^4 / \gamma_{\text{ax}}^2 + (s^2/a^2) \mathbb{E}[\tilde{\delta}^2] + 2(s/a) \text{Cov} + O(s^4/\tau^2)$.

Proof. $\hat{x}_0 = (x - s\varepsilon^*)/a = x(1 - s^2/\gamma^2)/a = (a\sigma_{\text{ax}}^2/\gamma^2)x$, so $\text{Var} = a^2 \sigma_{\text{ax}}^4/\gamma^2$; Tweedie's formula identifies this with $\mathbb{E}[X_0 | X_{\hat{t}}]$ [14]. The error term is the jump map applied to $\varepsilon^* - \tilde{\delta}$; the curvature term is the posterior-mean chord bias. $\square$

Consequences: strongly singular axes ($\sigma_{\text{ax}} \ll s$) are collapsed onto the manifold (the matcher must top up); fat axes ($\sigma_{\text{ax}} \gg s$) are untouched while injection fattens them (the matcher must contract); axes with $\sigma_{\text{ax}} \sim s$ sharpen by an axis-dependent factor strictly inside $(0, 1)$. Hence the required correction is signed and anisotropic: any correct matcher must be bidirectional and per-axis. (*We observe:* scalar matching leaves thin ratios at 0.62 to 0.65 on MNIST latents; per-axis matching achieves $|\text{thin} - 1| \leq 0.012$.)

E The residual channel (*we model*)

(*Role revised by measurement:* the floor anatomy of Section 5 identifies the single-pass remainder beyond the envelope as one measured channel — in-band profile sag, the ensemble-displacement effect having been absorbed into the hypothesis by (3) — with the residual beyond it small and slightly negative; this appendix's stochastic bound is retained as a consistency check on the residual field's magnitude, not as the remainder's explanation.)

The residual of the estimator about its affine $L^2(p_t)$ fit, $\delta := \hat{\varepsilon} - \hat{\kappa}_{\text{aff}}x - b$ (used in this appendix only, as an independent characterisation of the residual field; the floor proof does not rely on it), has measured band risk $r_{\text{res}}^2(t) \approx 0.02$ to 0.05 , flat in t and in sample size, with path-coherence time $\approx \nu t$, $\nu \approx 0.75$. Under a block-correlation model (coherent within blocks of length νt , independent across blocks, no systematic over-contraction), the injected variance obeys $\Phi_0(t_0) \geq c\nu\rho_0(\sigma^2 + t_0)$: decreasing in t_0 until $t_0 \sim \sigma^2$, then saturating, the same shape as the measured excess (knee at $t_0 \approx \sigma^2$, n -flat), with the single constant effective to within a factor ~ 2 . In this paper the channel plays a supporting role only: it is the measured $\sim 30\%$ single-pass remainder on top of the closed-form floor, which does not depend on it.

F Experimental details

Networks and training. 3-layer MLPs, width 256, SiLU, 32-dim sinusoidal time embedding; Adam, learning rate 2×10^{-3} ; batch 512 (ring) / 256 (MNIST); denoising score matching with $t \sim U(10^{-3}, 8)$; fresh network each generation.

Sampling. Reverse Euler-Maruyama on a 200-point geometric grid from $T = 8$; the ablation zeroes the $\sqrt{|dt|}\xi$ term only. Jump: $\hat{t} = 0.05$.

Anchor. Ring: radial (normal) bidirectional match of the pool’s radial variance to the stale reference’s, contract-by-scaling / top-up-by-noise. MNIST: per-point k NN-PCA ($k = 96$, $k_{\text{dim}} = 12$) charts from the reference; normal coordinates rescaled to the reference’s local normal eigenvalues in the same frames. Reference: $m = 2,000$ real samples drawn once at $g = 0$, never refreshed (provenance-free: the matcher is applied to the whole pool without labels).

Measured network profile. The fitted slope $\hat{\kappa}(t)$ is extracted by affine regression of $\hat{\varepsilon}$ against x on fiber lines at each t ; it tracks $\kappa(t)$ (0.95 to 0.97 for $t \geq 0.1$), saturates at $\bar{\kappa} \approx 3.9$ inside the band, and never exceeds κ (n -independent across 1k to 16k) — a data-width statement: nets trained on fat pools ($\sqrt{w} \gtrsim 0.10$) show mild overshoot at mid times, so the class-(N) constant applies at the data width while class (U) needs no such assumption. The decomposition identity $(\kappa - \hat{\kappa})^2 \gamma^2 + r_{\text{res}}^2 = r_{\text{tot}}^2$ checks numerically along the path.

t_0 -dependence of the response slope. The probe’s net slope $c \in \{0.00, \dots, 0.13\}$ is measured at its sampling point $t_0 = 5 \times 10^{-3}$. At the recursion’s operating point $t_0 = 10^{-4}$, the directly measured per-pass map over the same widths has $dv/dw \approx 0.68$ ($c \approx -0.32$): once the sampler traverses the whole band, the deficit channel dominates and e^2 decreases over this range — still within the law’s requirement ($e^{2'} > -1$ keeps T monotone and concave, and only strengthens contraction), and consistent with the plateau measurements ($e^2(w^*) = 3.9\text{--}5.0 \sigma^2$ falling with the fixed point’s width). The $\lambda = 0.02$ runaway then implies the slope turns back positive in the far-fat regime ($w \gg 10 \sigma^2$), outside the anatomy’s measured range; Assumption 2’s global nondecreasing-concave form is a regime model, exact where identified.

Anatomy of c . Training single-width tubes at $w \in \{0.0009, \dots, 0.0169\}$ and measuring $dv/dw - 1$ gives net $c \in \{0.00, 0.05, 0.13, 0.10, 0.05\}$ (small, positive, concave) as the sum of a deficit response ≈ -0.68 (a fatter pool needs a smaller peak slope) and an injection response $\approx +0.73$ (a fatter pool damps injected variance less). The same probe’s per-width slope fits give the capacity-degradation fit $\bar{\kappa}(w) = 4.43 - 11.6\sqrt{w}$ used in Section 5, which refines this anatomy into three channels: writing $\frac{d}{dw}\Phi_{\text{det}}(w, \bar{\kappa}(w)) = \partial_w \Phi_{\text{det}}|_{\bar{\kappa}} + (\partial_{\bar{\kappa}} \Phi_{\text{det}}) \bar{\kappa}'(w)$, the mid-range values are a deficit ≈ -0.94 and a capacity-degradation term $\approx +0.7$ (the latter is not a fit artifact: raw centered differences of the five measured points, with no functional form at all, give $+0.78/ +0.68/ +0.67$ at the three interior widths vs. the fit’s $+0.84/ +0.70/ +0.65$). The derived deterministic response is therefore ≈ -0.23 mid-range, and the measured total ($+0.10$ to $+0.13$) leaves a residual attributed to injection of $\approx +0.3$ to $+0.4$. Because injection is *defined* as this residual, the three channels sum to the measurement by construction; the non-trivial content is that the residual comes out stable, positive, and of the size the residual-channel model (Appendix E) expects, and that most of what the two-channel anatomy above attributes to injection is really $\bar{\kappa}$ degrading on fatter pools. We trust the split mid-range only: at the thinnest width $\bar{\kappa}'(w)$ from the $\sqrt{w}$ fit diverges, an artifact of the parametrization, not a mechanism.